

# Neural networks & Neural language models

CS 4804: Introduction to AI

*Fall 2025*

<https://tuvllms.github.io/ai-fall-2025/>

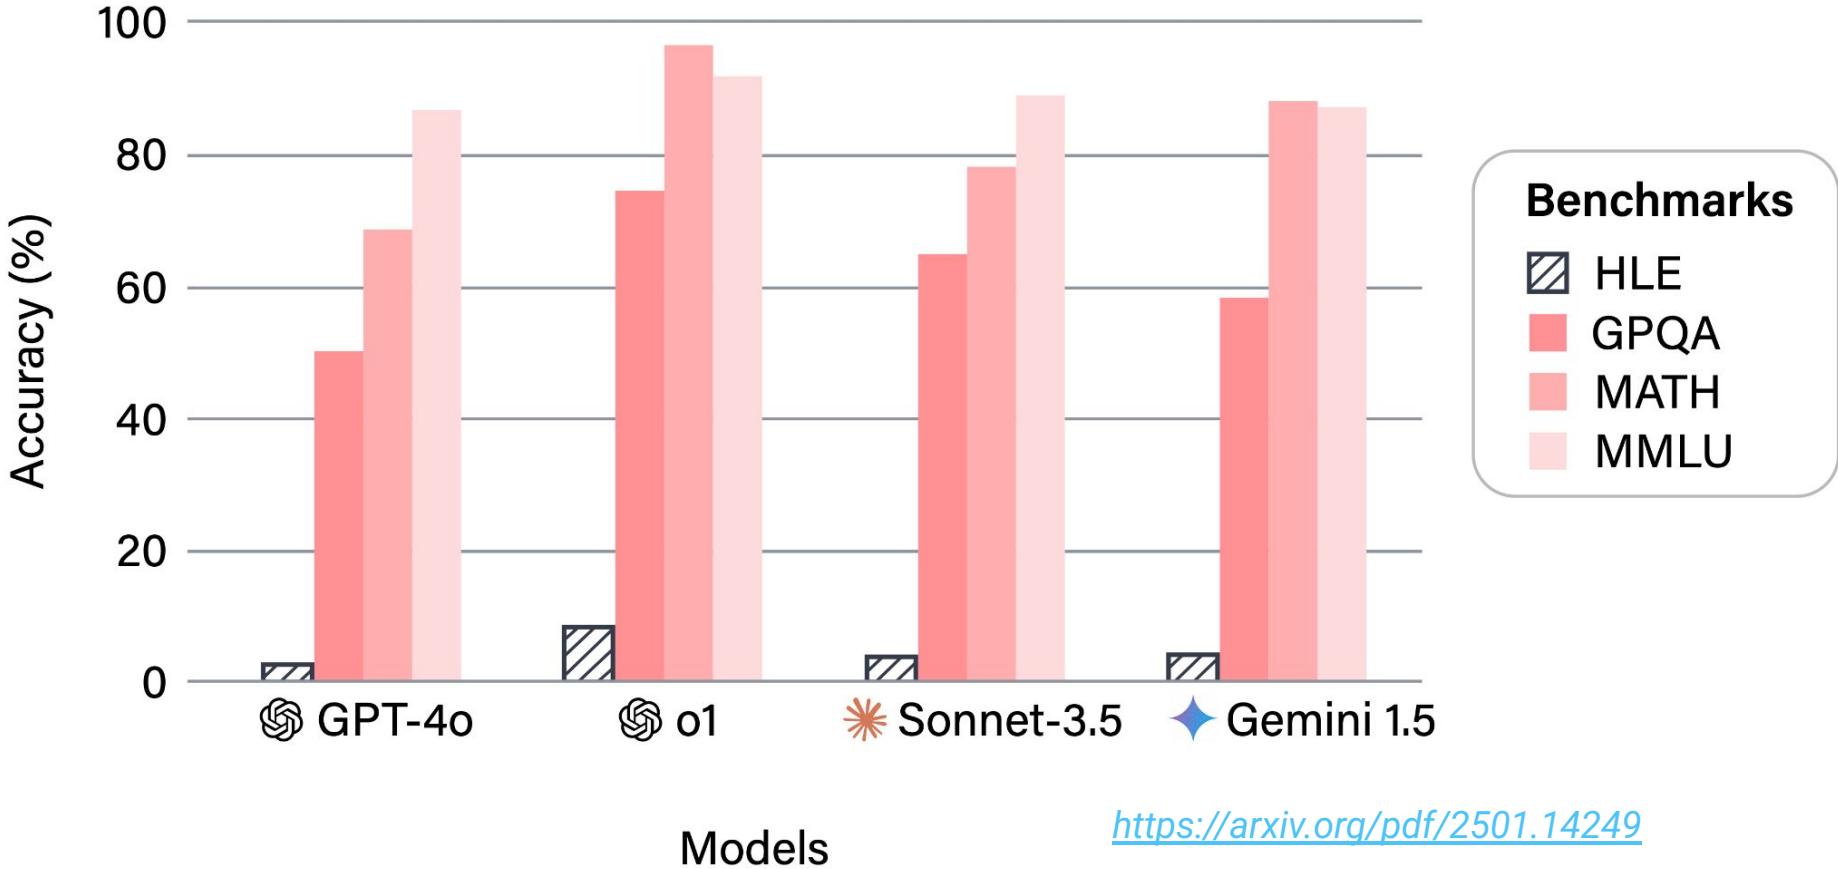
Tu Vu



# Logistics

- Office hours starting this week
  - both in-person and via Zoom (**links available on Piazza**)
- Homework 0 released (**due September 16<sup>th</sup>**)
- Final project group
  - Search for teammates on Piazza  
<https://piazza.com/class/meqiibrwtql168/post/5> or  
reach out to us at [cs4804instructors@gmail.com](mailto:cs4804instructors@gmail.com)
  - Google form for submitting group information available on Piazza (**due September 5<sup>th</sup>**)

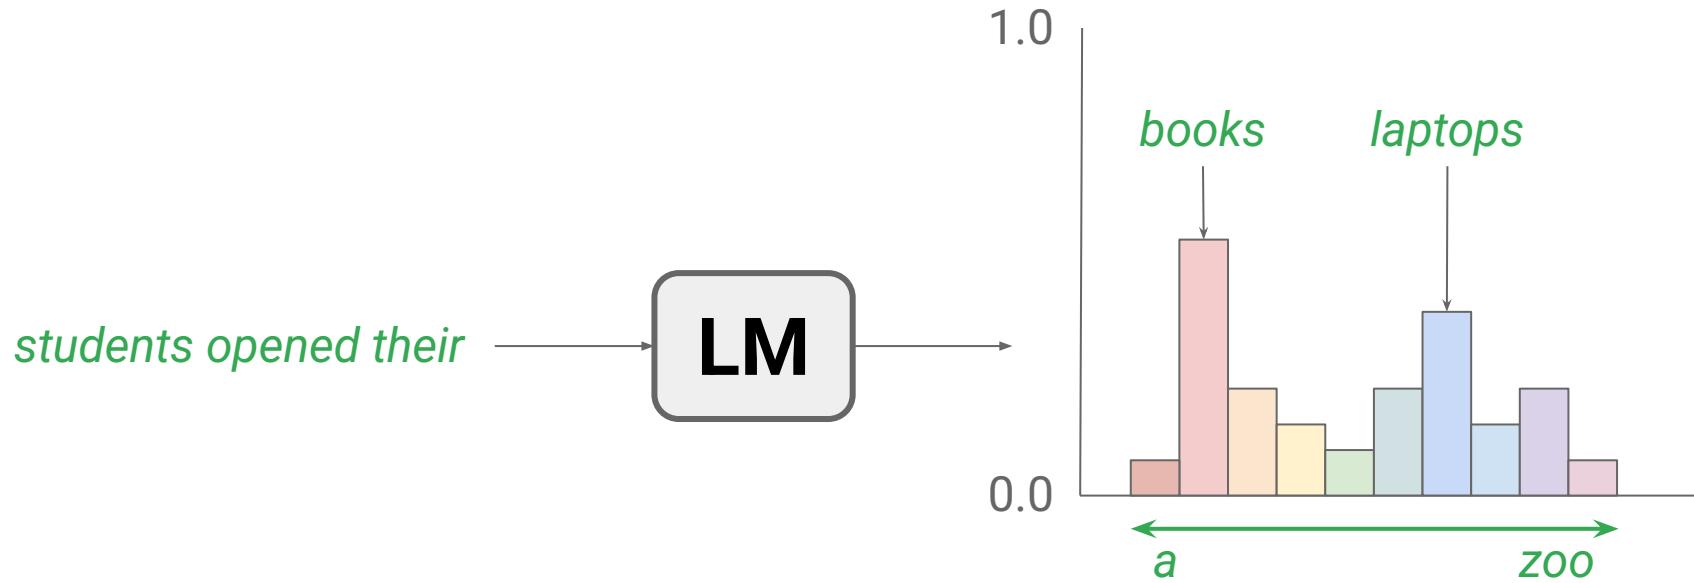
# Humanity's Last Exam



<https://arxiv.org/pdf/2501.14249>

# Language modeling review

- Predict the next word, or a *probability distribution* over possible next words



# Language modeling review (cont'd)

- Language models
  - compute

$$P(w_1, w_2, \dots, w_n)$$

or

$$P(w_j | w_1, w_2, \dots, w_{j-1})$$

$$P(w_1, w_2, \dots, w_n) = P(w_1) \times P(w_2 | w_1) \times \dots \times P(w_n | w_1, w_2, \dots, w_{n-1})$$

# N-gram language models review

- Use maximum likelihood estimation (MLE)

$$P(\text{"laptops"} \mid \text{"students opened their"}) = \frac{\text{Count}(\text{"students opened their laptops"})}{\text{Count}(\text{"students opened their"})}$$

# Perplexity

$$\text{perplexity}(W) = P(w_1 w_2 \dots w_N)^{-\frac{1}{N}}$$

We normalize by the number of words N by taking the Nth root

$$= \sqrt[N]{\frac{1}{P(w_1 w_2 \dots w_N)}}$$

Or we can use the chain rule to expand the probability of  $W$ :

$$\text{perplexity}(W) = \sqrt[N]{\prod_{i=1}^N \frac{1}{P(w_i | w_1 \dots w_{i-1})}}$$

# Perplexity as Weighted Average Branching Factor

- Suppose a sentence consists of random digits.  
What is the perplexity of this sentence for a model that assigns a probability of 1/10 to each digit?

$$\text{PP}(W) = P(w_1 w_2 \dots w_N)^{-\frac{1}{N}}$$

$$= \left(\frac{1}{10^N}\right)^{-\frac{1}{N}}$$

$$= \frac{1}{10}^{-1}$$

$$= 10$$

**Given any prefix, how many next words does the model consider reasonable?**

# In practice, we use log probs

$$\log \prod p(w_i | w_{i-1}) = \sum \log p(w_i | w_{i-1})$$

logs to avoid  
numerical underflow

**sentence:** I love love love love love the movie

$$p(\text{i}) \cdot p(\text{love})^5 \cdot p(\text{the}) \cdot p(\text{movie}) = 5.95374181\text{e-}7$$

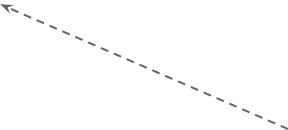
$$\log p(\text{i}) + 5 \log p(\text{love}) + \log p(\text{the}) + \log p(\text{movie})$$

$$= -14.3340757538$$

source: Mohit Iyyer

## In practice, we use log probs (cont'd)

$$\text{perplexity}(W) = \exp\left(-\frac{1}{N} \sum_i^N \log p(w_i | w_{<i})\right)$$



perplexity is the  
exponentiated token-level  
negative log-likelihood

# Problems with n-gram language models

$P(\text{"laptops"} \mid \text{"students opened their"}) =$

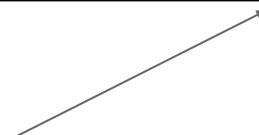
$$\frac{\text{Count}(\text{"students opened their laptops"})}{\text{Count}(\text{"students opened their"})}$$

**What if “*students opened their laptops*” never occurred in training data?**

# Problems with n-gram language models (cont'd)

	<b>i</b>	<b>want</b>	<b>to</b>	<b>eat</b>	<b>chinese</b>	<b>food</b>	<b>lunch</b>	<b>spend</b>
<b>i</b>	0.002	0.33	0	0.0036	0	0	0	0.00079
<b>want</b>	0.0022	0	0.66	0.0011	0.0065	0.0065	0.0054	0.0011
<b>to</b>	0.00083	0	0.0017	0.28	0.00083	0	0.0025	0.087
<b>eat</b>	0	0	0.0027	0	0.021	0.0027	0.056	0
<b>chinese</b>	0.0063	0	0	0	0	0.52	0.0063	0
<b>food</b>	0.014	0	0.014	0	0.00092	0.0037	0	0
<b>lunch</b>	0.0059	0	0	0	0	0.0029	0	0
<b>spend</b>	0.0036	0	0.0036	0	0	0	0	0

Need to store  $V^n$  counts for an n-gram model!



# Problems with n-gram language models (cont'd)

- Treat semantically similar prefixes independently of each other

*"students opened their \_\_"*

*"pupils opened their \_\_"*

*"scholars opened their \_\_"*

*"students began reading their \_\_"*

**Shouldn't we share  
information across  
these prefixes?**

# Matrix-vector multiplication

Matrix  $A$  (dimensions  $4 \times 3$ ):

$$A = \begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \\ a_{41} & a_{42} & a_{43} \end{bmatrix}$$

Vector  $x$  (dimensions  $3 \times 1$ ):

$$x = \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}$$

Resulting vector  $b$  (dimensions  $4 \times 1$ ):

$$b = A \cdot x = \begin{bmatrix} a_{11}x_1 + a_{12}x_2 + a_{13}x_3 \\ a_{21}x_1 + a_{22}x_2 + a_{23}x_3 \\ a_{31}x_1 + a_{32}x_2 + a_{33}x_3 \\ a_{41}x_1 + a_{42}x_2 + a_{43}x_3 \end{bmatrix}$$

# Softmax function

For a vector  $y = [y_1, y_2, \dots, y_V]$  of dimension  $V$ , the softmax transformation is calculated as:

$$\text{softmax}(y) = \left[ \frac{e^{y_1}}{\sum e^y}, \frac{e^{y_2}}{\sum e^y}, \dots, \frac{e^{y_V}}{\sum e^y} \right]$$

where  $\sum e^y = e^{y_1} + e^{y_2} + \dots + e^{y_V}$ .

# Word representations / embeddings

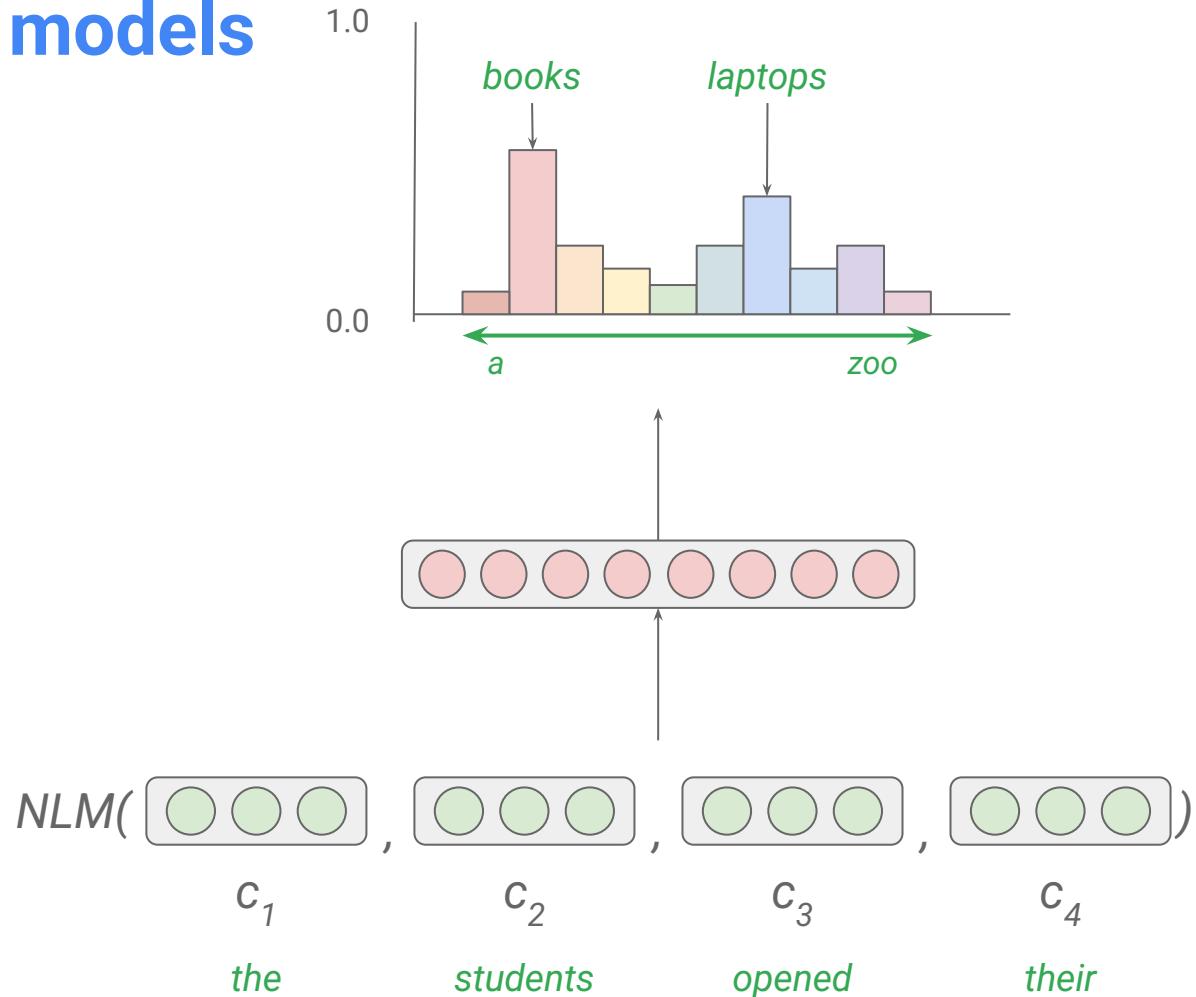
- High-dimensional / sparse / one-hot representations
- Low-dimensional / dense representations

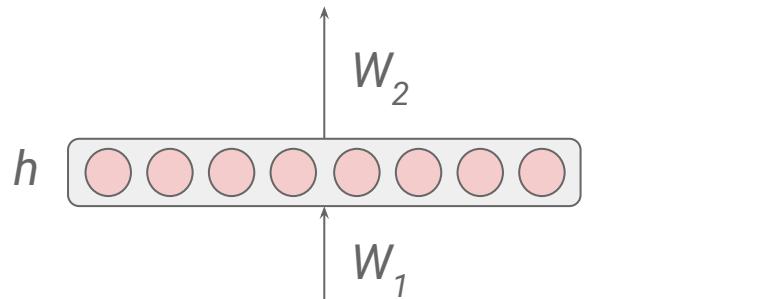
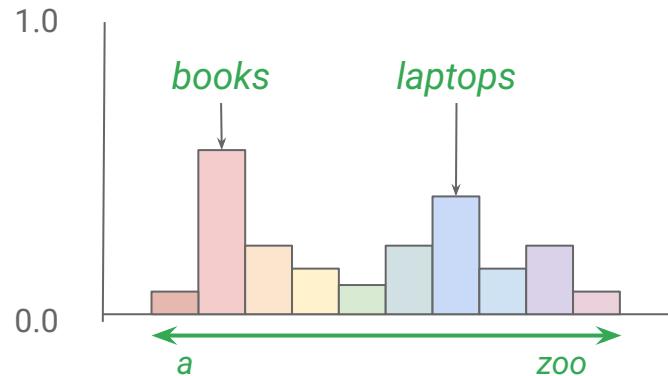
# Word representations / embeddings (cont'd)

```
▶ #What is the vector representation for a word?  
w2v_model['computer']
```

```
array([ 1.07421875e-01, -2.01171875e-01,  1.23046875e-01,  2.11914062e-01,  
       -9.13085938e-02,  2.16796875e-01, -1.31835938e-01,  8.30078125e-02,  
       2.02148438e-01,  4.78515625e-02,  3.66210938e-02, -2.45361328e-02,  
       2.39257812e-02, -1.60156250e-01, -2.61230469e-02,  9.71679688e-02,  
      -6.34765625e-02,  1.84570312e-01,  1.70898438e-01, -1.63085938e-01,  
      -1.09375000e-01,  1.49414062e-01, -4.65393066e-04,  9.61914062e-02,  
      1.68945312e-01,  2.60925293e-03,  8.93554688e-02,  6.49414062e-02,  
      3.56445312e-02, -6.93359375e-02, -1.46484375e-01, -1.21093750e-01,  
     -2.27539062e-01,  2.45361328e-02, -1.24511719e-01, -3.18359375e-01,  
     -2.20703125e-01,  1.30859375e-01,  3.66210938e-02, -3.63769531e-02,  
     -1.13281250e-01,  1.95312500e-01,  9.76562500e-02,  1.26953125e-01,  
      6.59179688e-02,  6.93359375e-02,  1.02539062e-02,  1.75781250e-01,  
     -1.68945312e-01,  1.21307373e-03, -2.98828125e-01, -1.15234375e-01,  
      5.66406250e-02, -1.77734375e-01, -2.08984375e-01,  1.76757812e-01,  
      2.38037109e-02, -2.57812500e-01, -4.46777344e-02,  1.88476562e-01,  
      5.51757812e-02,  5.02929688e-02, -1.06933594e-01,  1.89453125e-01,  
     -1.16210938e-01,  8.49609375e-02, -1.71875000e-01,  2.45117188e-01,  
     -1.73828125e-01, -8.30078125e-03,  4.56542969e-02, -1.61132812e-02,  
      1.86523438e-01, -6.05468750e-02, -4.17480469e-02,  1.82617188e-01,  
     2.20703125e-01, -1.22558594e-01, -2.55126953e-02, -3.08593750e-01,  
      9.13085938e-02,  1.60156250e-01,  1.70898438e-01,  1.19628906e-01,  
      7.08007812e-02, -2.64892578e-02, -3.08837891e-02,  4.06250000e-01,  
     -1.01562500e-01,  5.71289062e-02, -7.26318359e-03, -9.17968750e-02,  
     -1.50390625e-01, -2.55859375e-01,  2.16796875e-01, -3.63769531e-02,  
     2.24609375e-01.  8.00781250e-02.  1.56250000e-01.  5.27343750e-02.
```

# Neural language models





**output distribution**

$$\hat{y} = \text{softmax}(W_2 h)$$

$x$        $c_1$        $c_2$        $c_3$        $c_4$   
*the*      *students*      *opened*      *their*

# Composition functions

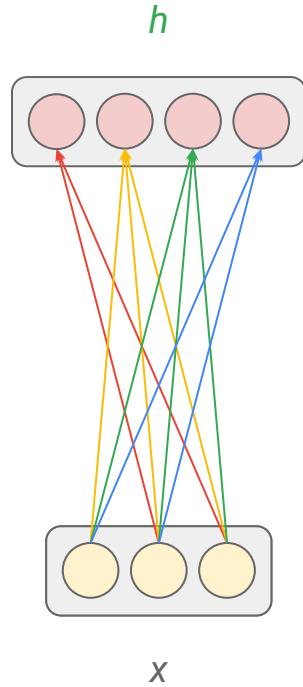
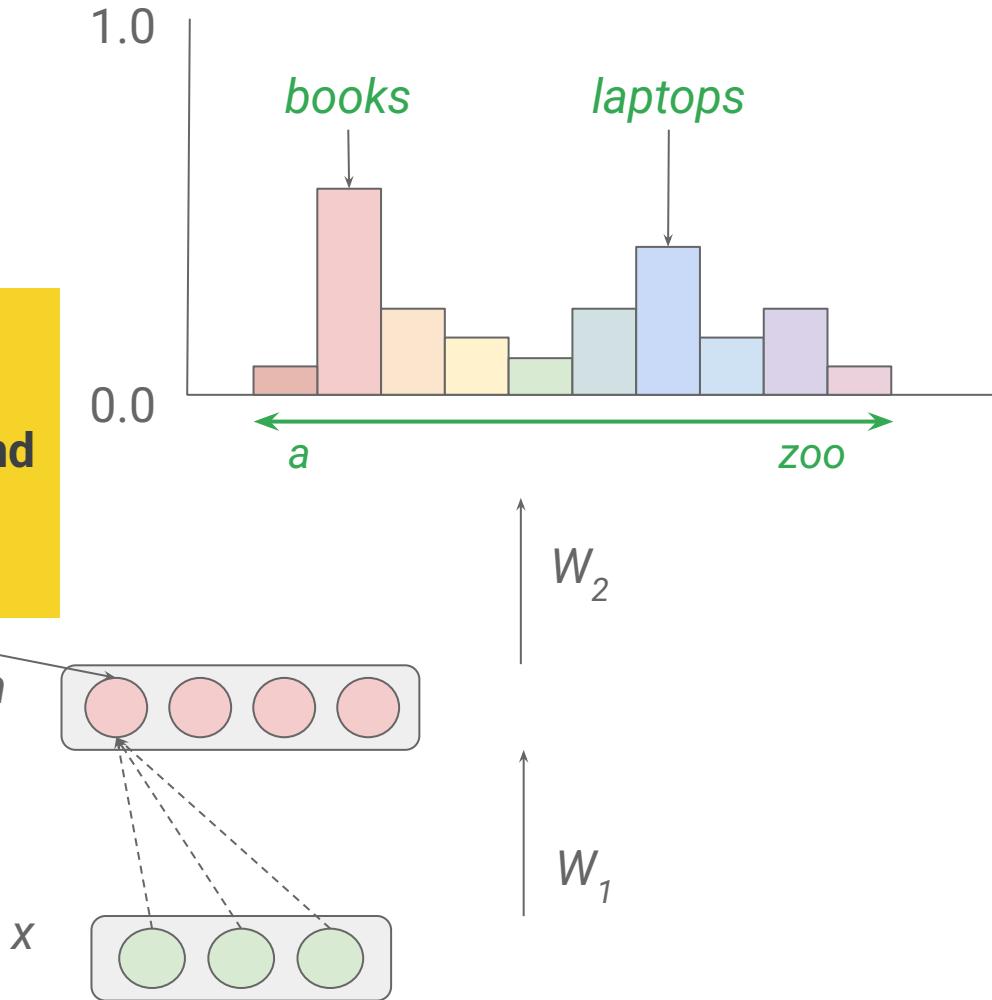
- Element-wise functions
  - e.g., just sum up all of the word embeddings
- Concatenation
- Feedforward neural networks
- Convolutional neural networks
- Recurrent neural networks
- Transformers

# Feedforward neural language model

## hidden layer

$$h = f(W_1 x)$$

hidden unit:  
taking a weighted  
sum of its inputs and  
then applying a  
non-linearity



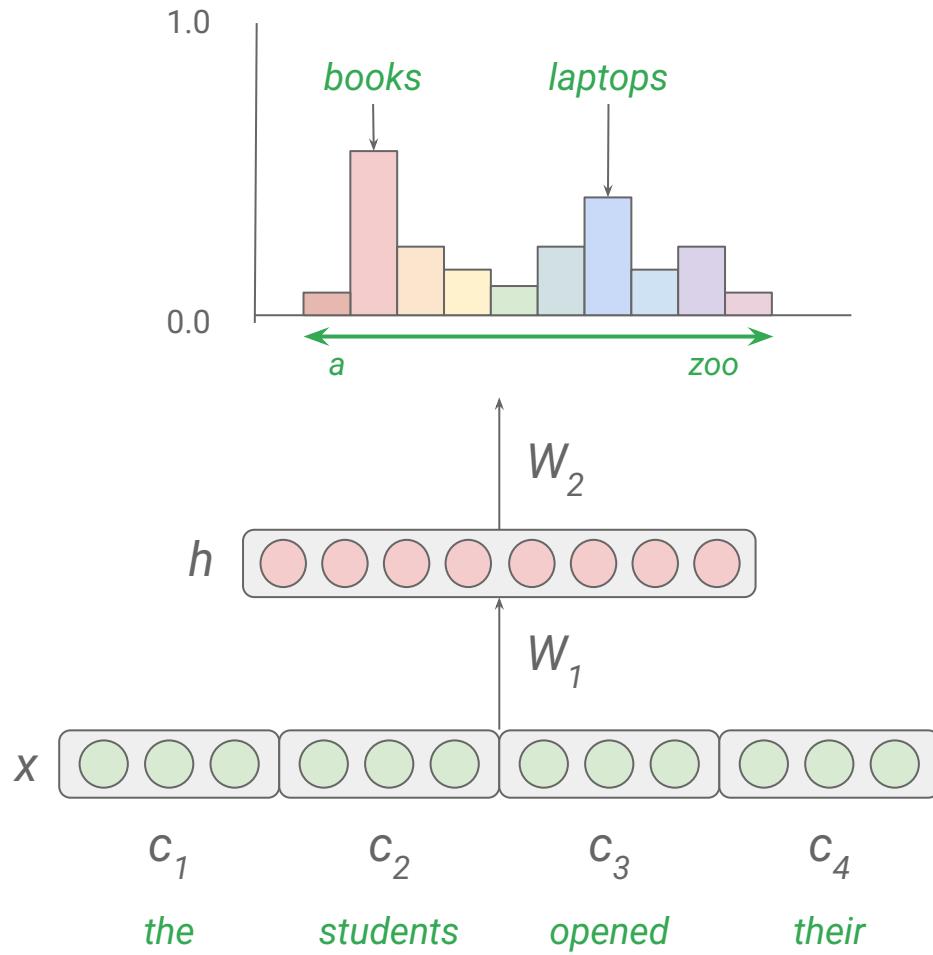
## hidden layer

$$h = f(W_1 x)$$

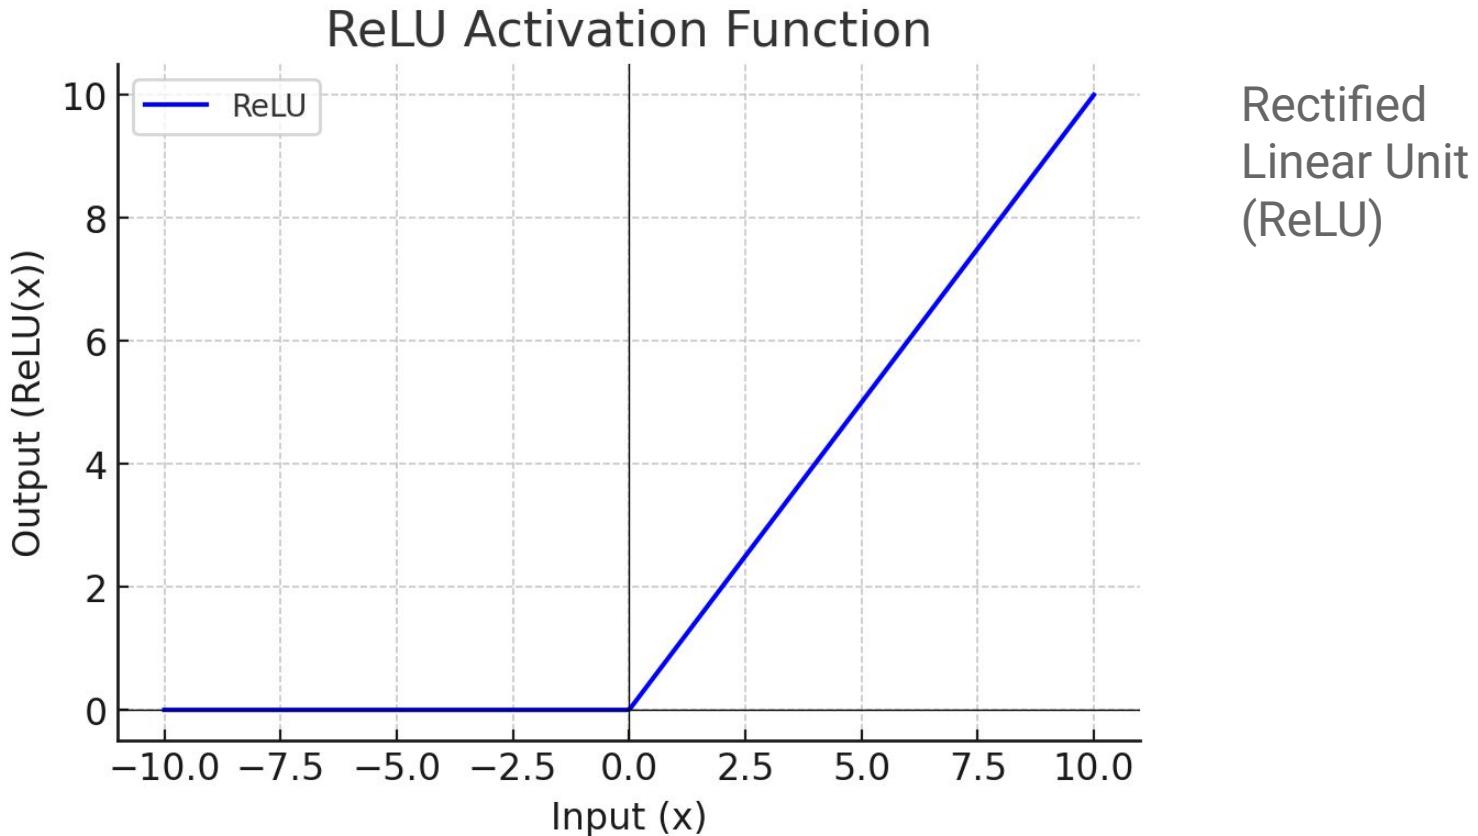
f is a non-linear activation function to model non-linear relationships between words

## output distribution

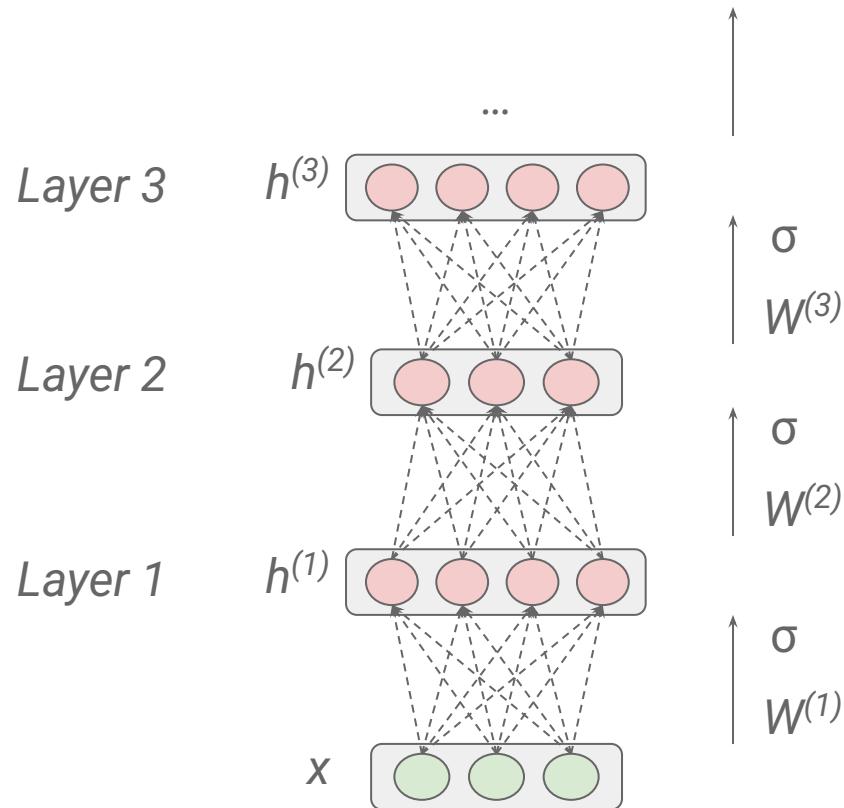
$$\hat{y} = \text{softmax}(W_2 h)$$



# Activation functions



# Deep neural networks

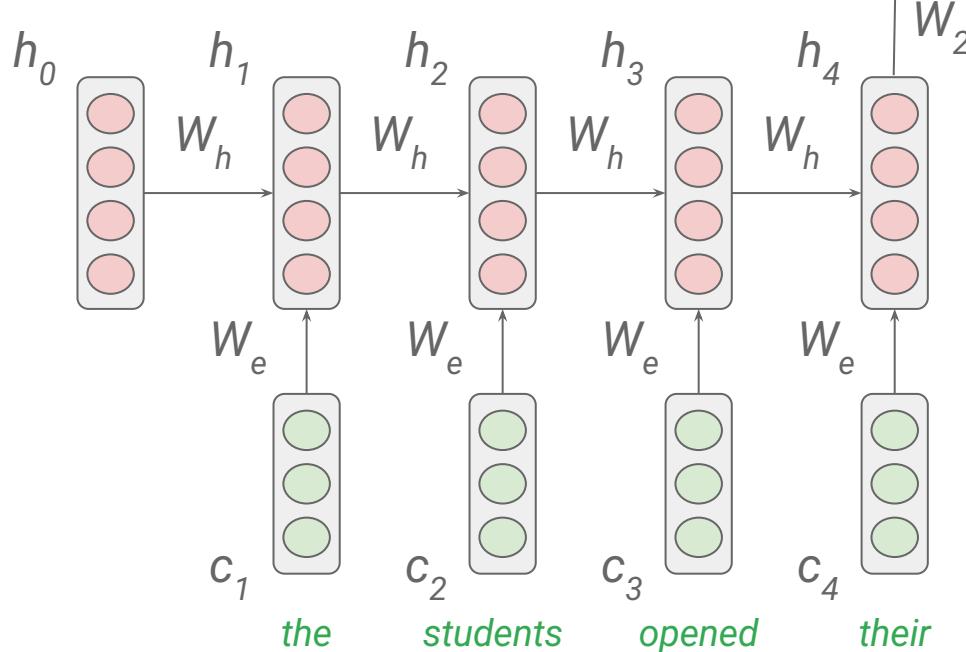
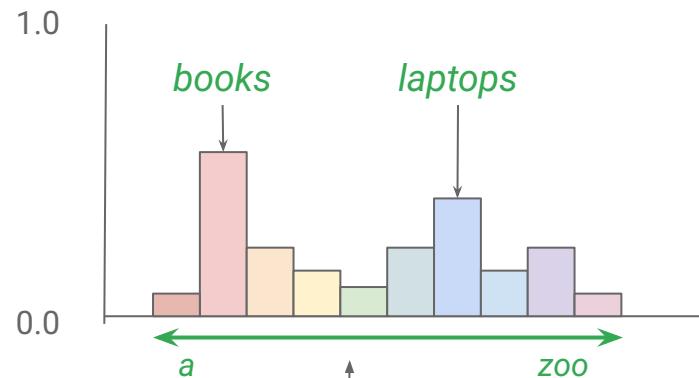


hierarchical representations, where each layer builds upon the previous one

# Recurrent neural networks (RNNs)

## hidden states

$$h^{(t)} = f(W_h h^{(t-1)} + W_e c^t)$$



## output distribution

$$\hat{y} = \text{softmax}(W_2 h^{(n-1)})$$

# Recurrent neural networks (RNNs)

- RNNs advantages
  - can handle much longer histories
  - can generalize better over contexts of similar words
  - are more accurate at word-prediction
- RNNs disadvantages
  - are much more complex
  - are slower and need more energy to train
  - and are less interpretable than n-gram models

**Thank you!**