

# Mixture of Experts

**CS 4804: Introduction to AI**  
*Fall 2025*

<https://tuvllms.github.io/ai-fall-2025/>

**Tu Vu**



# Logistics

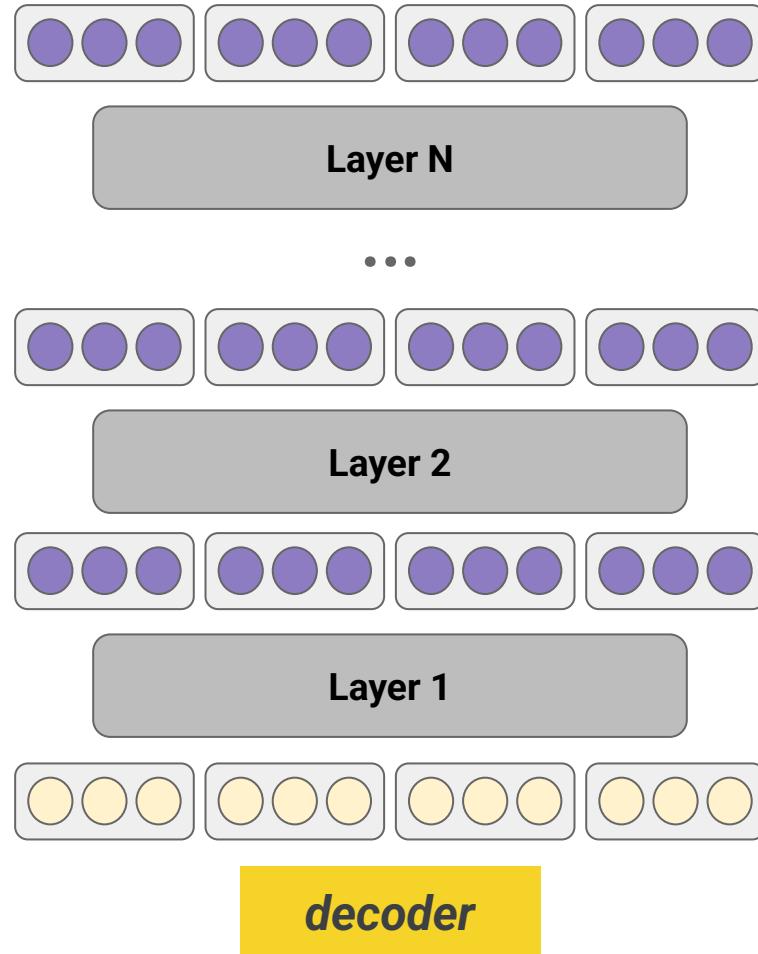
- No class next Tuesday 10/28 (Tu traveling), resume as usual next Thursday
- We are sending feedback for final project proposals
  - Please follow the template
- Quiz 2 released due 10/30
- HW 2 will be released later today due 11/13
- Teaching & learning evaluation: 11/4
- Final presentations: 12/4 & 12/9

# AI Browsers

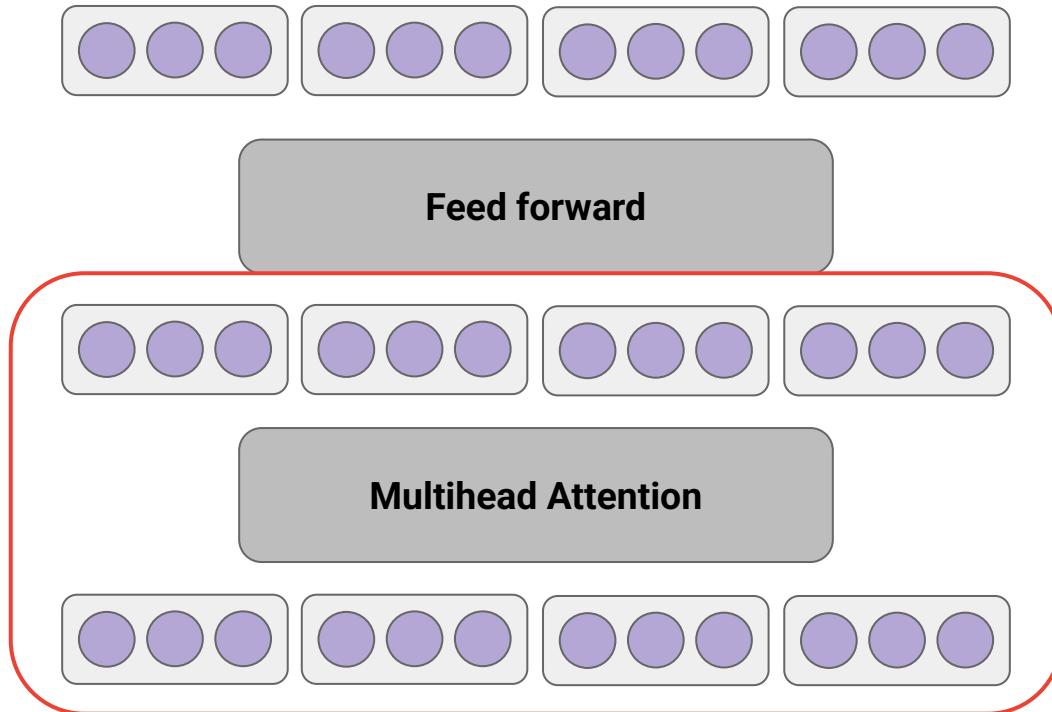
- OpenAI's Atlas
  - <https://openai.com/index/introducing-chatgpt-atlas/>
- Perplexity's Comet
  - <https://www.perplexity.ai/comet>

# Decoder-only Transformer review

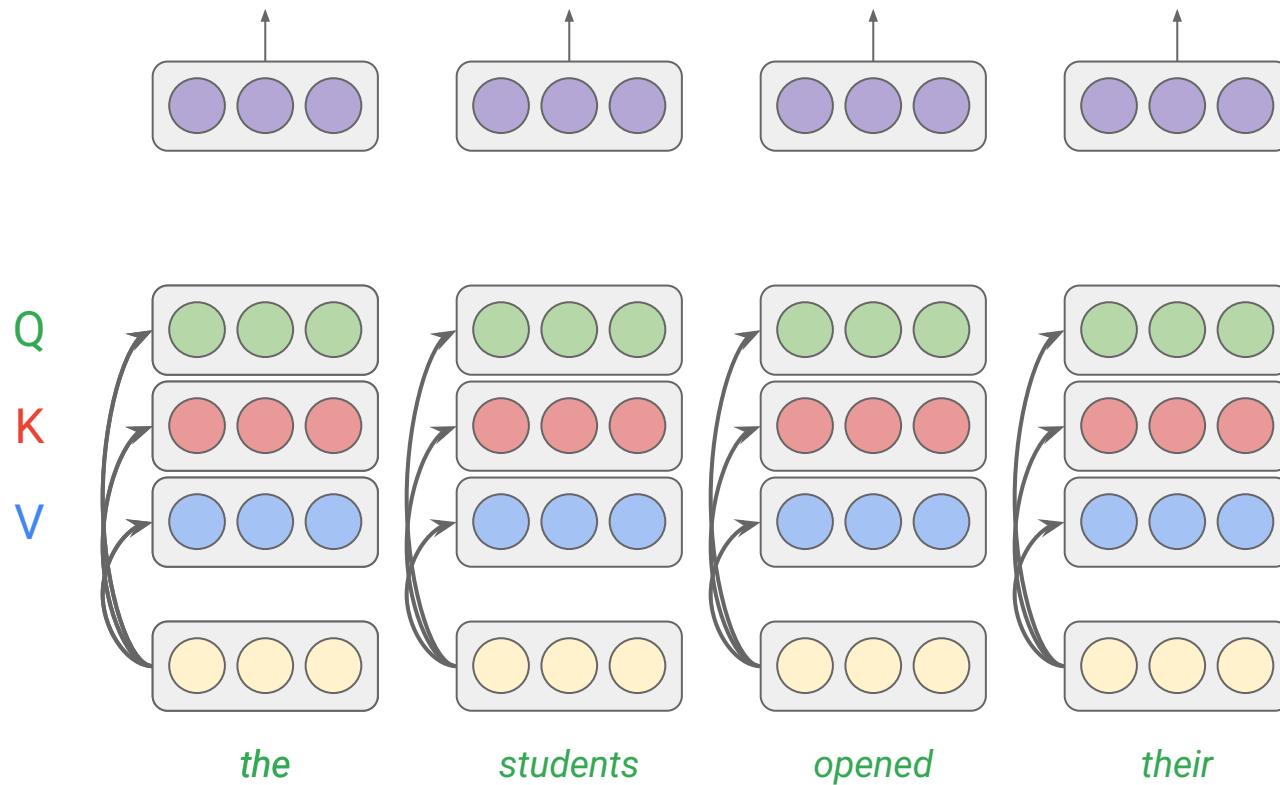
# Transformer (N layers)



# Transformer decoder



# Attention



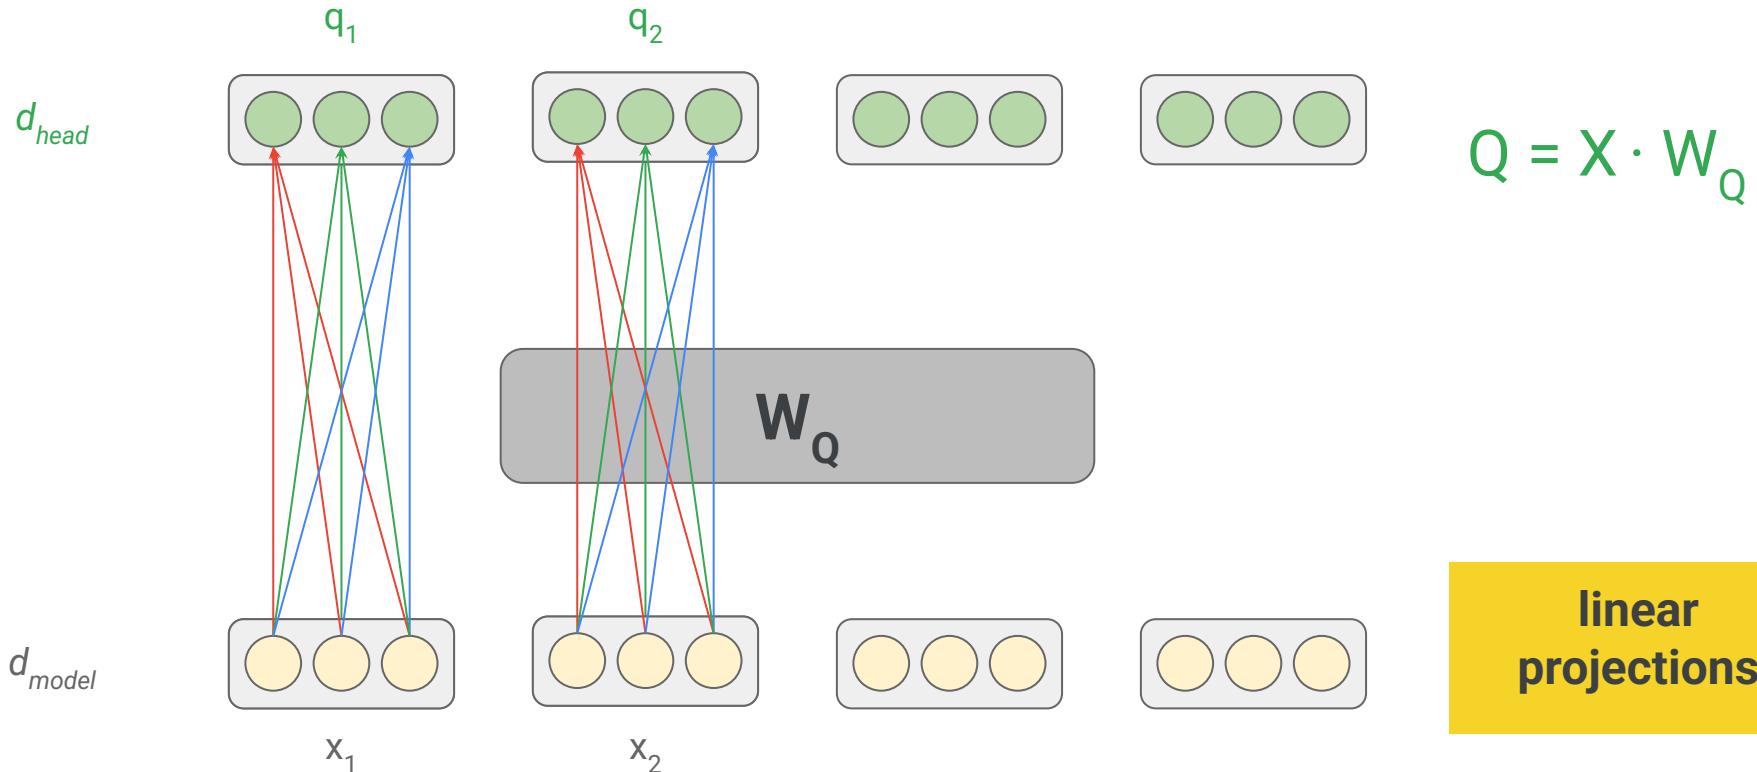
$$Q = X \cdot W_Q$$

$$K = X \cdot W_K$$

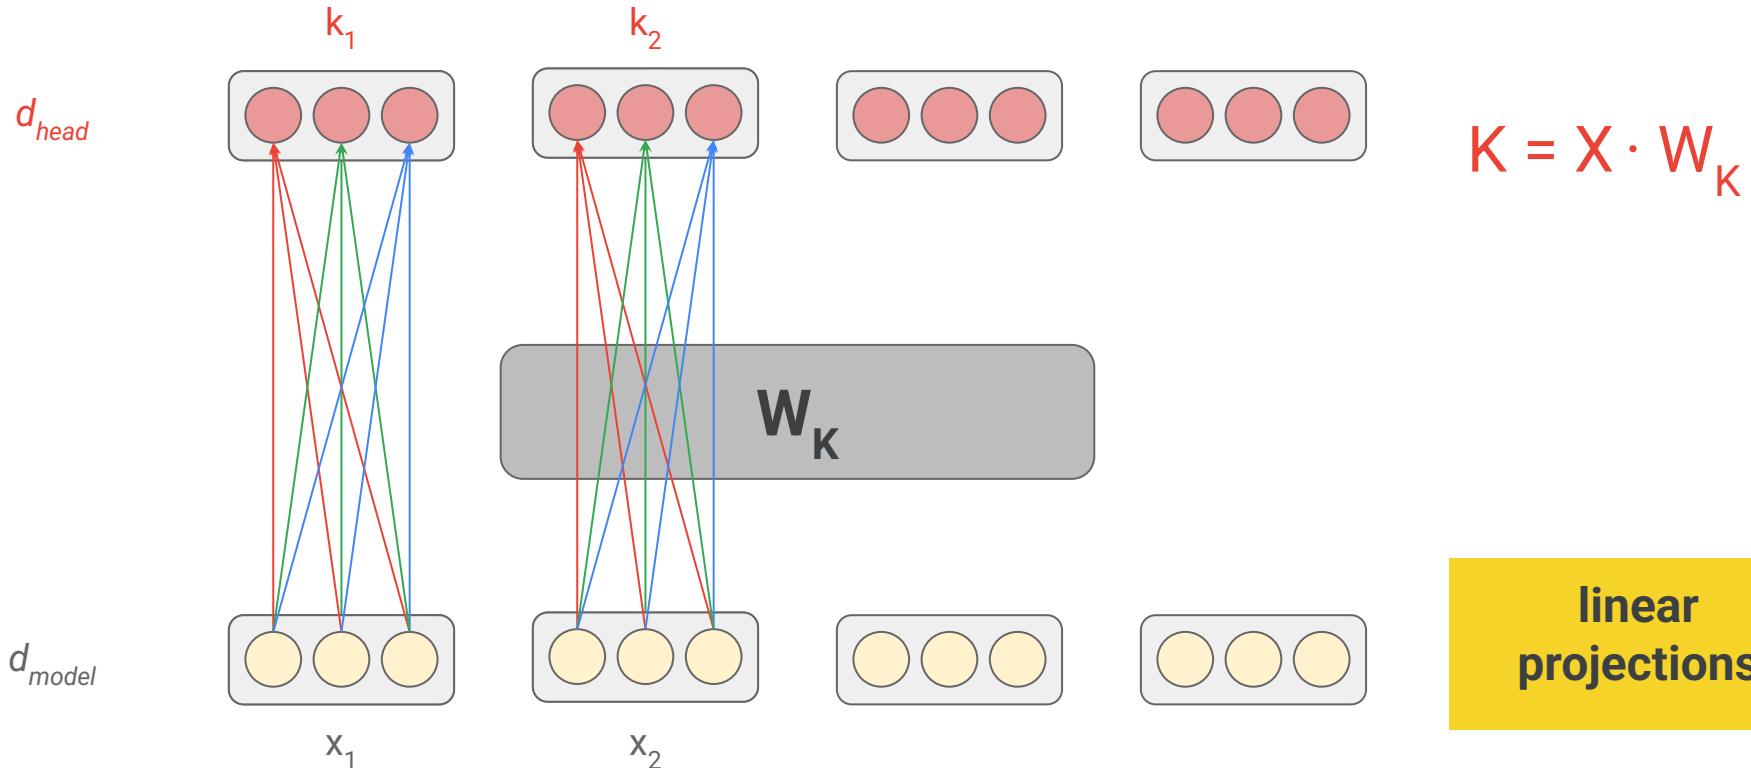
$$V = X \cdot W_V$$

linear  
projections

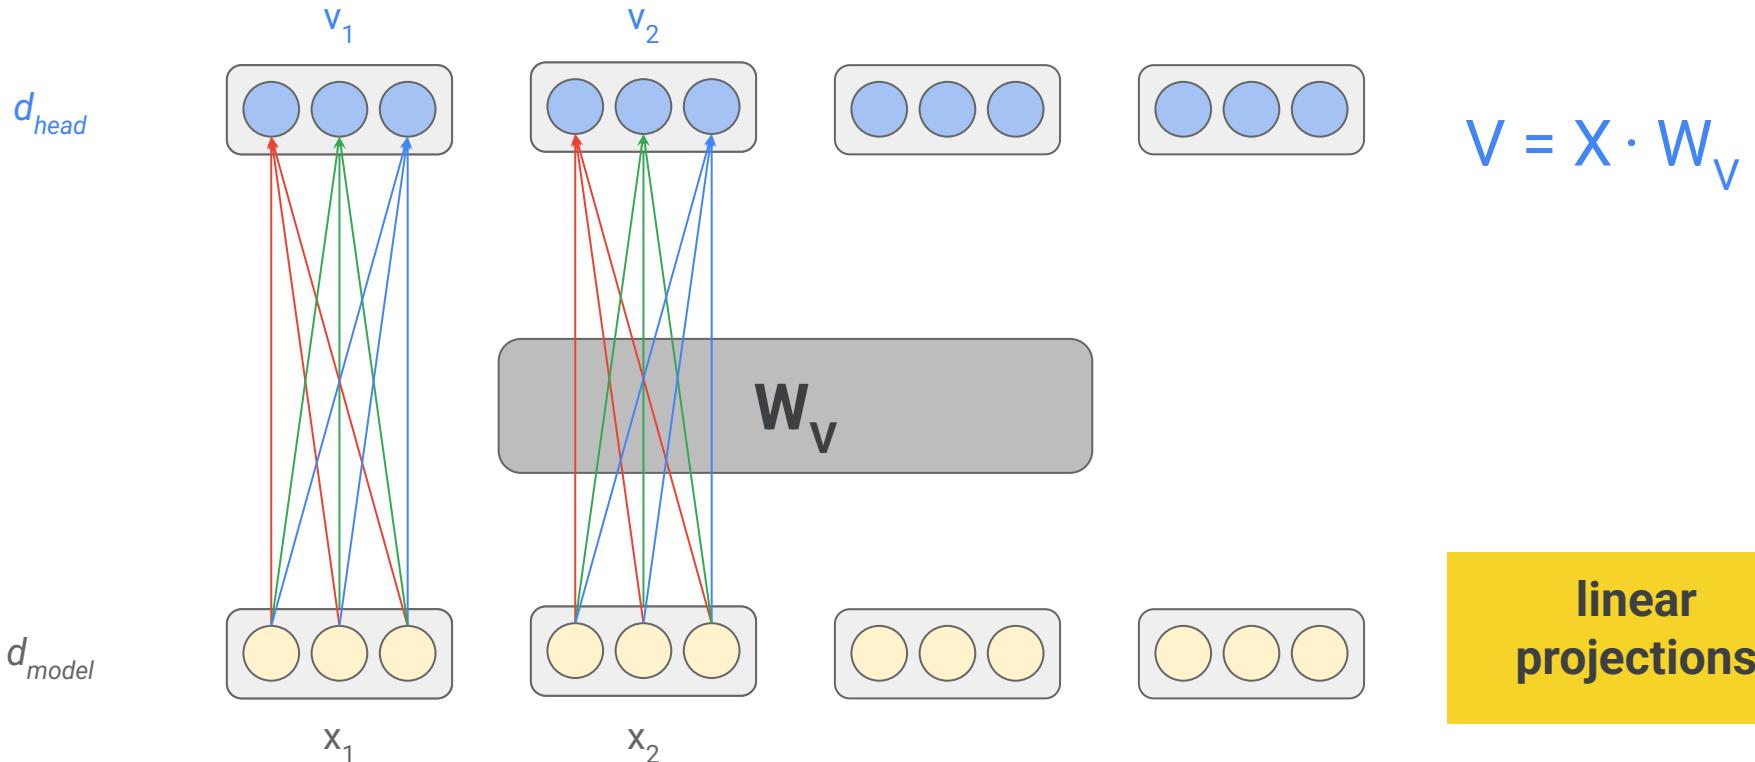
# Query vectors



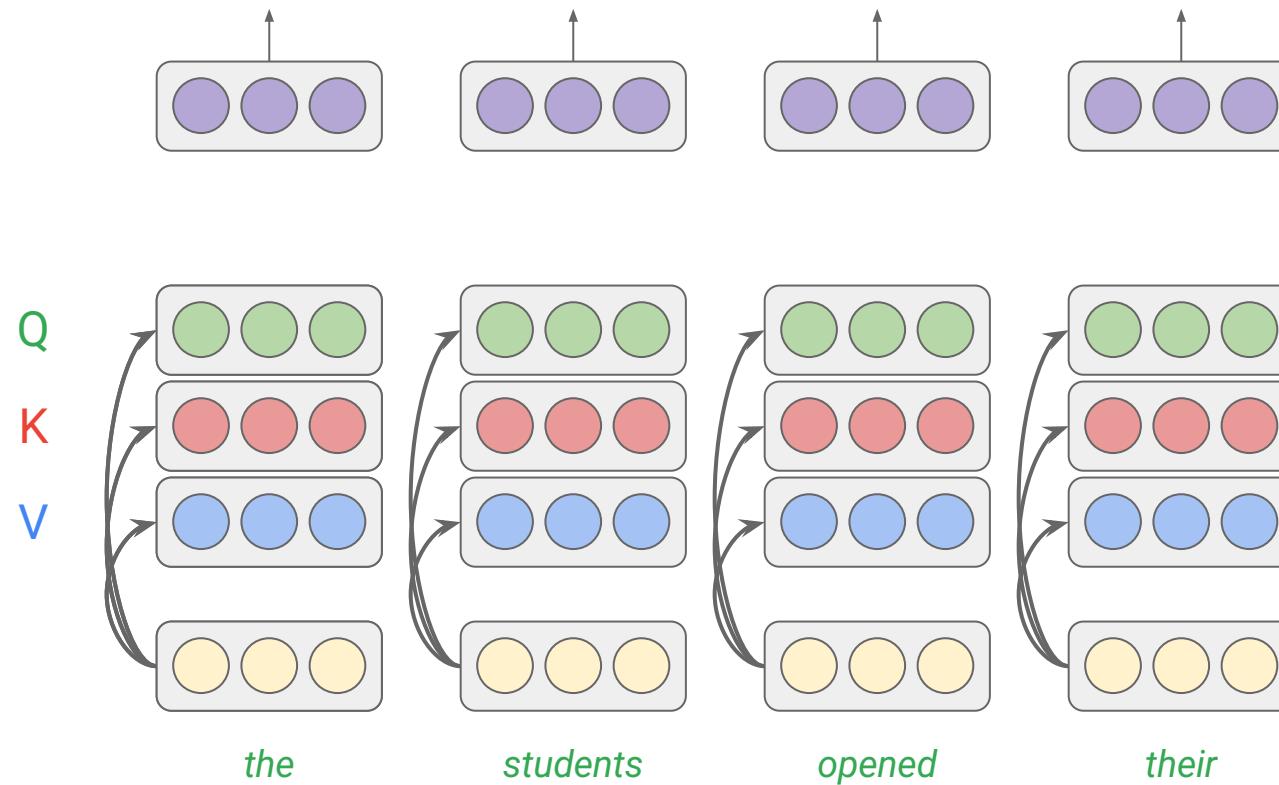
# Key vectors



# Value vectors



# Attention (cont'd)



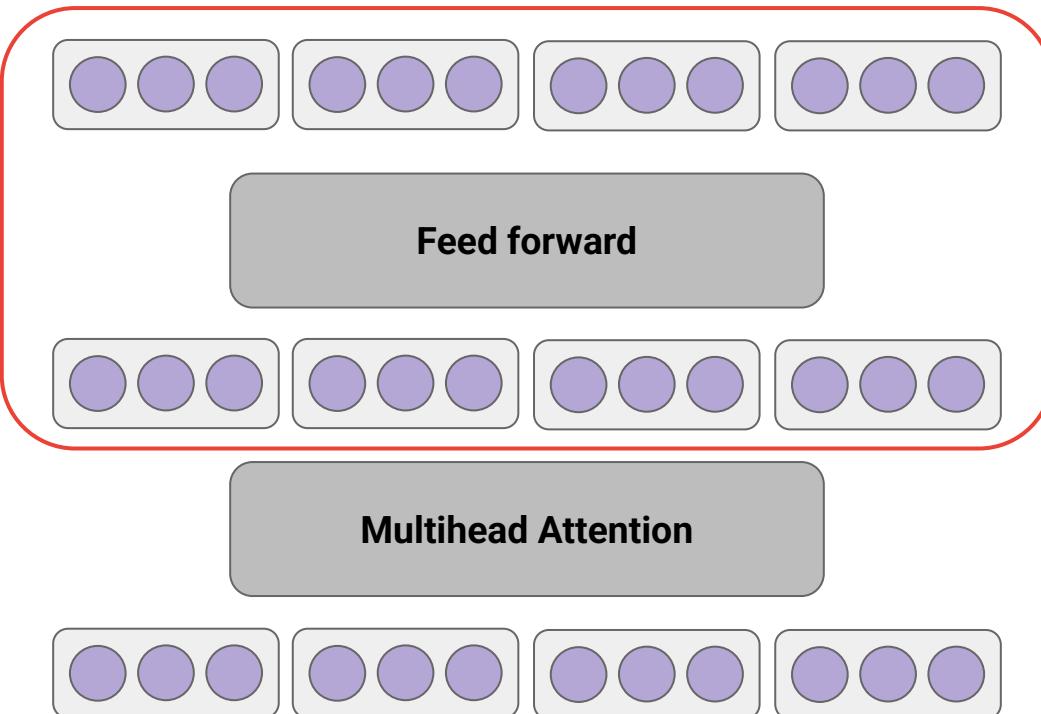
$$Q = X \cdot W_Q$$

$$K = X \cdot W_K$$

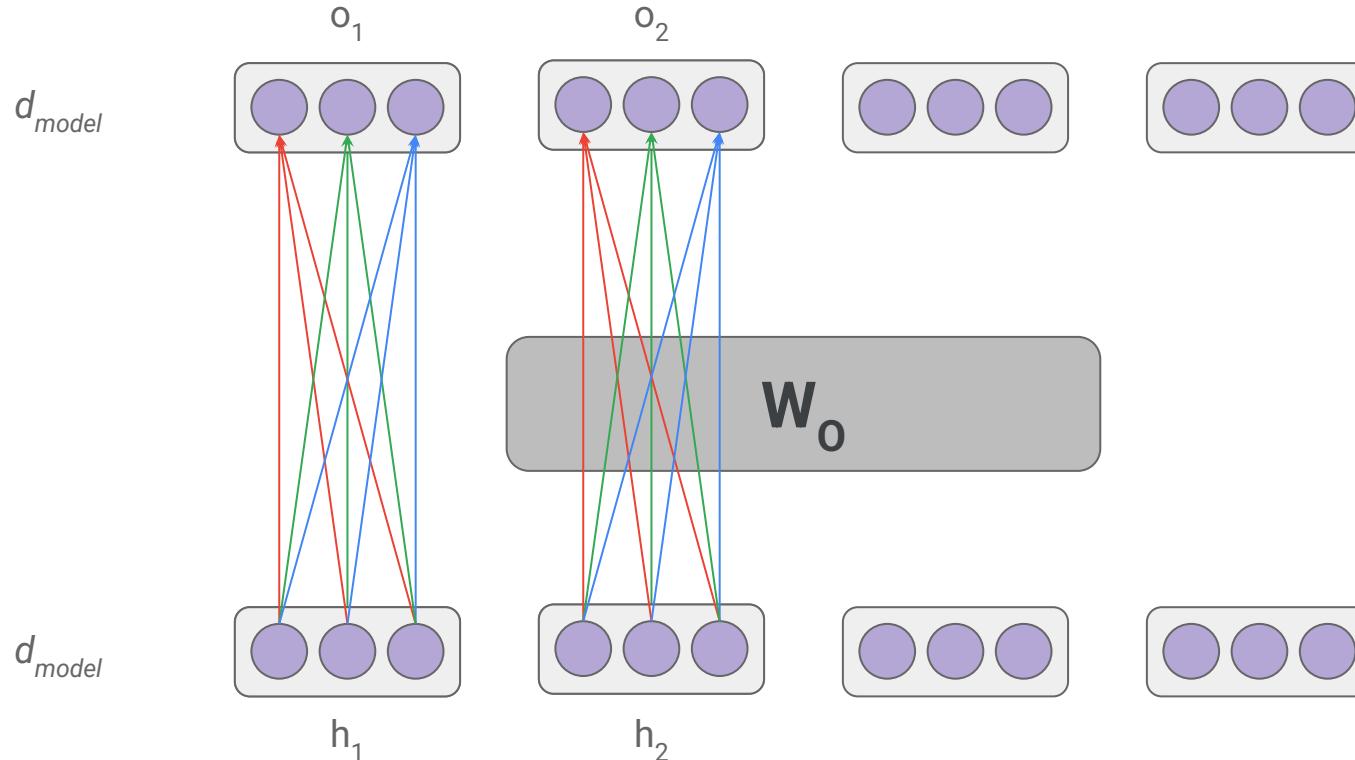
$$V = X \cdot W_V$$

linear  
projections

# Transformer decoder



# output vectors



$$O = H \cdot W_o$$

linear  
projections

# What are mixture-of-experts?

- The gpt-oss models are autoregressive Mixture-of-Experts (MoE) transformers
  - **gpt-oss-120b:** 36 layers (116.8B total parameters and 5.1B “active” parameters per token)
  - **gpt-oss-20b:** 24 layers (20.9B total and 3.6B active parameters)
- Llama 4

## ※ Llama 4

Llama 4 release

∞ [meta-llama/Llama-4-Scout-17B-16E-Instruct](#)

Image-Text-to-Text • ∴ 109B • Updated May 22 • ↓ 182k • ⚡ • ❤ 1.12k

∞ [meta-llama/Llama-4-Scout-17B-16E](#)

Image-Text-to-Text • ∴ 109B • Updated Apr 9 • ↓ 14.4k • ❤ 207

∞ [meta-llama/Llama-4-Maverick-17B-128E-Instruct](#)

Image-Text-to-Text • ∴ 402B • Updated May 22 • ↓ 20.6k • ⚡ • ❤ 417

∞ [meta-llama/Llama-4-Maverick-17B-128E-Instruct-FP8](#)

Image-Text-to-Text • ∴ 402B • Updated May 22 • ↓ 193k • ⚡ • ❤ 137

# Llama 4: Leading Multimodal Intelligence

Newest model suite offering unrivaled speed and efficiency

## Llama 4 Behemoth

**288B** active parameter, **16** experts

**2T** total parameters

The most intelligent teacher model for distillation

Preview

## Llama 4 Maverick

**17B** active parameters, **128** experts

**400B** total parameters

Native multimodal with **1M** context length

Available

## Llama 4 Scout

**17B** active parameters, **16** experts

**109B** total parameters

Industry leading **10M** context length  
Optimized inference

Available

# The Bitter Lesson

“The biggest lesson that can be read from 70 years of AI research is that general methods that leverage computation are ultimately the most effective, and by a large margin.”

Rich Sutton, 2019

Simple architectures—backed by a generous computational budget, data set size and parameter count—surpass more complicated algorithms

# OUTRAGEOUSLY LARGE NEURAL NETWORKS: THE SPARSELY-GATED MIXTURE-OF-EXPERTS LAYER

Noam Shazeer<sup>1</sup>, Azalia Mirhoseini<sup>\*1</sup>, Krzysztof Maziarz<sup>\*2</sup>, Andy Davis<sup>1</sup>, Quoc Le<sup>1</sup>, Geoffrey Hinton<sup>1</sup> and Jeff Dean<sup>1</sup>

<sup>1</sup>Google Brain, {noam,azalia,andydavis,qvl,geoffhinton,jeff}@google.com

<sup>2</sup>Jagiellonian University, Cracow, krzysztof.maziarz@student.uj.edu.pl

## ABSTRACT

The capacity of a neural network to absorb information is limited by its number of parameters. Conditional computation, where parts of the network are active on a per-example basis, has been proposed in theory as a way of dramatically increasing model capacity without a proportional increase in computation. In practice, however, there are significant algorithmic and performance challenges. In this work, we address these challenges and finally realize the promise of conditional



TECHNOLOGY | ARTIFICIAL INTELLIGENCE

## Google Paid \$2.7 Billion to Bring Back an AI Genius Who Quit in Frustration



Noam Shazeer

[Google](#)

Verified email at google.com

[Deep Learning](#)



Get access to  
**10M+ Characters**

Sign up in just ten seconds

### TITLE

### CITED BY    YEAR

[Attention is all you need](#)

209038    2017

A Vaswani, N Shazeer, N Parmar, J Uszkoreit, L Jones, AN Gomez, ...

Advances in neural information processing systems 30

[Exploring the limits of transfer learning with a unified text-to-text transformer](#)

27346    2020

C Raffel, N Shazeer, A Roberts, K Lee, S Narang, M Matena, Y Zhou, W Li, ...

Journal of machine learning research 21 (140), 1-67

# Switch Transformers: Scaling to Trillion Parameter Models with Simple and Efficient Sparsity

**William Fedus\***

LIAMFEDUS@GOOGLE.COM

**Barret Zoph\***

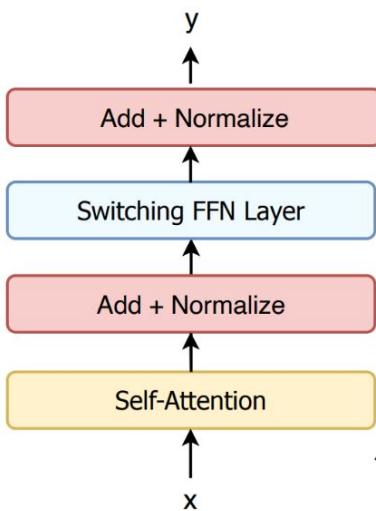
BARRETZOPH@GOOGLE.COM

**Noam Shazeer**

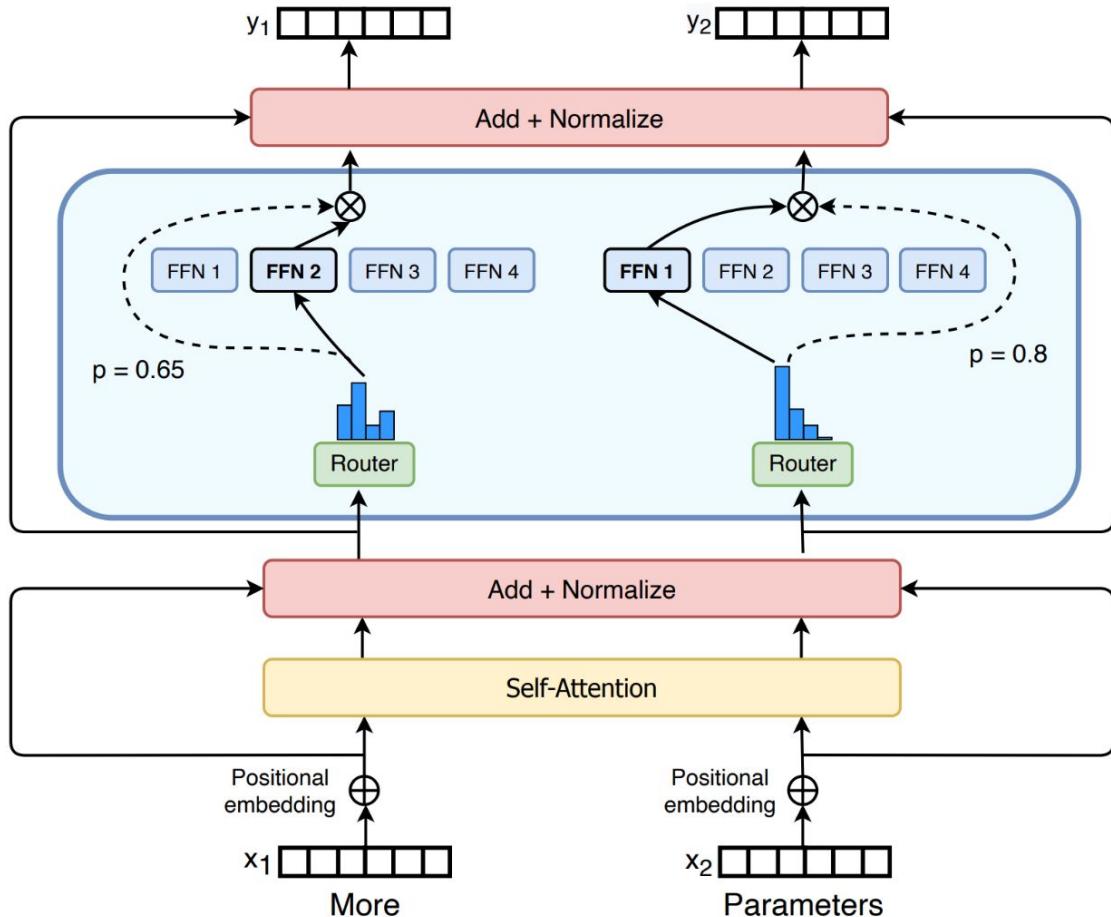
NOAM@GOOGLE.COM

*Google, Mountain View, CA 94043, USA*

# Switch Transformers



*The layer operates independently on the tokens*



# Switch Transformers

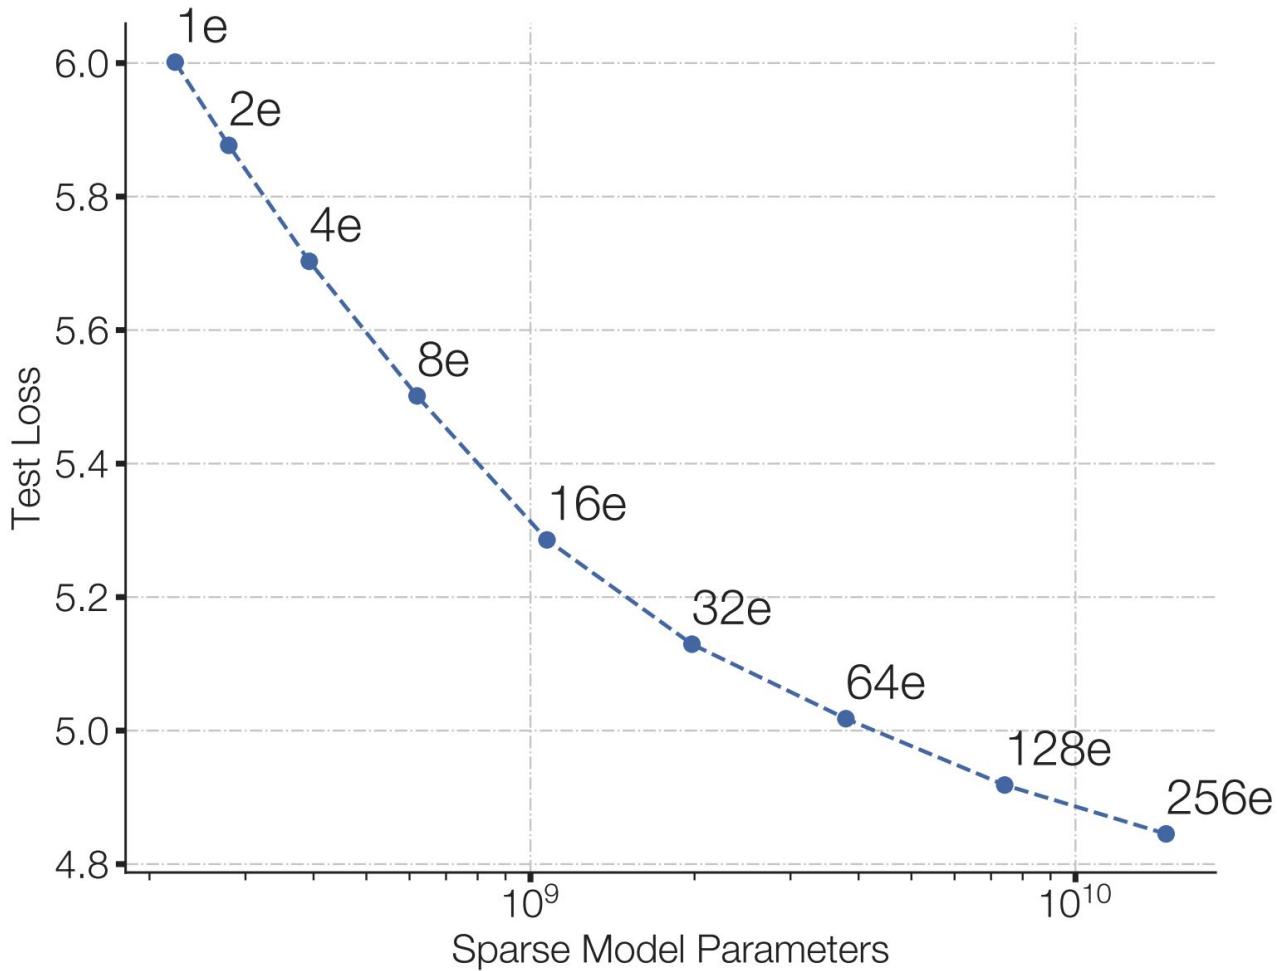
- Vanilla Transformer
  - densely-activated
- Switch Transformer
  - sparsely-activated expert model
  - with an outrageous number of parameters—but a constant computational cost (!)
  - pretraining up to *trillion* parameter models and achieving a 4x speedup over the T5-XXL (11B)

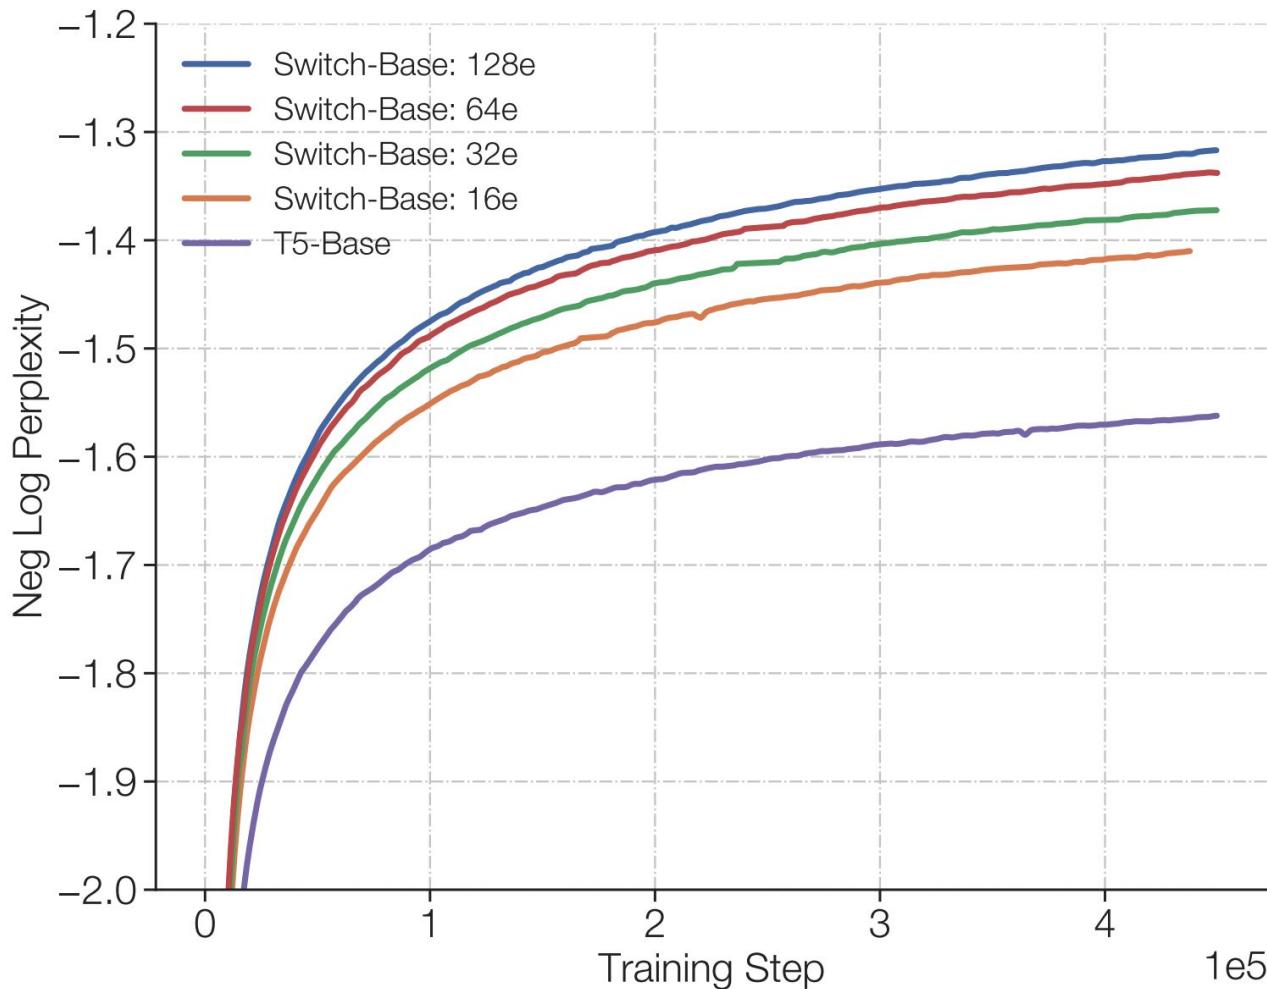
# Rethinking Mixture-of-Experts

- Shazeer et al. (2017)
  - Each token is processed by every expert, and their outputs are combined
  - This increases computational cost linearly with the number of experts, so adding experts makes training and inference more expensive
- Switch layer
  - Only a small subset of experts are activated per token
  - This allows the model to scale to many more experts while keeping the computational cost per token roughly constant

# Rethinking Mixture-of-Experts (con't)

- Shazeer et al. (2017)
  - routing to  $k > 1$  experts
  - intuition: learning to route would not work without the ability to compare at least two experts
- Switch layer
  - routes to only a *single* expert
  - preserves model quality
  - reduces routing computation
  - performs better





# Mixture of Expert Routing

The MoE layer takes as an input a token representation  $x$  and then routes this to the best determined top- $k$  experts, selected from a set  $\{E_i(x)\}_{i=1}^N$  of  $N$  experts.

The router variable  $W_r$  produces logits  $h(x) = W_r \cdot x$ , which are normalized via a softmax distribution over the available  $N$  experts at that layer. The gate value for expert  $i$  is given by:

$$p_i(x) = \frac{e^{h(x)_i}}{\sum_j e^{h(x)_j}}$$

The top- $k$  gate values are selected for routing the token  $x$ . If  $T$  is the set of selected top- $k$  indices, then the output computation of the layer is the linearly weighted combination of each expert's computation on the token by the gate value:

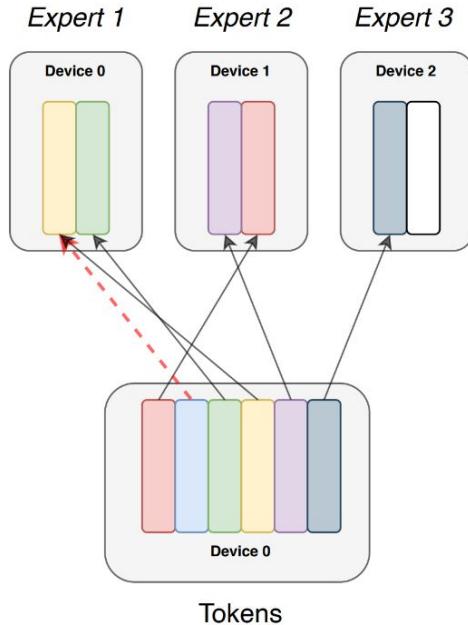
$$y = \sum_{i \in T} p_i(x) E_i(x)$$

$$\text{expert capacity} = \left( \frac{\text{tokens per batch}}{\text{number of experts}} \right) \times \text{capacity factor}$$

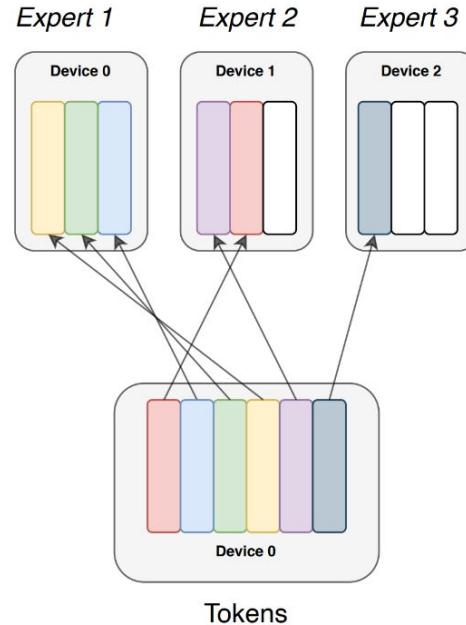
## Terminology

- Experts:** Split across devices, each having their own unique parameters. Perform standard feed-forward computation.
- Expert Capacity:** Batch size of each expert. Calculated as  $(\text{tokens\_per\_batch} / \text{num\_experts}) * \text{capacity\_factor}$
- Capacity Factor:** Used when calculating expert capacity. Expert capacity allows more buffer to help mitigate token overflow during routing.

(Capacity Factor: 1.0)



(Capacity Factor: 1.5)



Across Device Communication

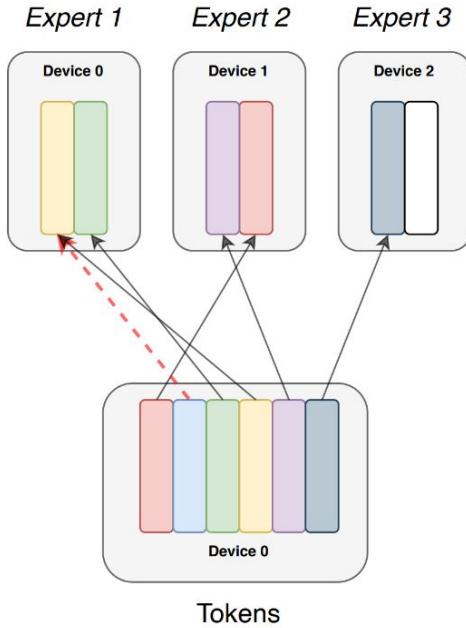
*What would happen with the blue (dropped) token?*

$$\text{expert capacity} = \left( \frac{\text{tokens per batch}}{\text{number of experts}} \right) \times \text{capacity factor}$$

## Terminology

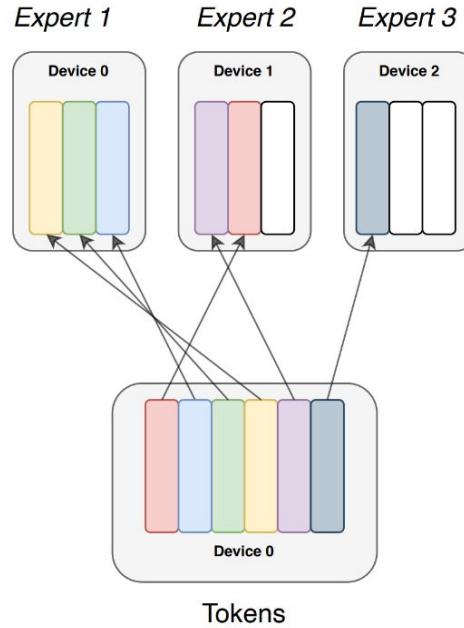
- Experts:** Split across devices, each having their own unique parameters. Perform standard feed-forward computation.
- Expert Capacity:** Batch size of each expert. Calculated as  $(\text{tokens\_per\_batch} / \text{num\_experts}) * \text{capacity\_factor}$
- Capacity Factor:** Used when calculating expert capacity. Expert capacity allows more buffer to help mitigate token overflow during routing.

(Capacity Factor: 1.0)



overflow

(Capacity Factor: 1.5)



wasted computation & memory

# An auxiliary load balancing loss

Given  $N$  experts indexed by  $i = 1$  to  $N$  and a batch  $B$  with  $T$  tokens, the auxiliary loss is computed as the scaled dot-product between vectors  $f$  and  $P$ :

$$\text{loss} = \alpha \cdot N \sum_{i=1}^N f_i \cdot P_i$$

$f_i$ : fraction of tokens to expert  $i$

- For each token in the batch, check which expert got chosen.
- Count how many times expert  $i$  was picked.
- Divide by the total number of tokens  $T$  to get the fraction.

$P_i$ : router probability for expert  $i$

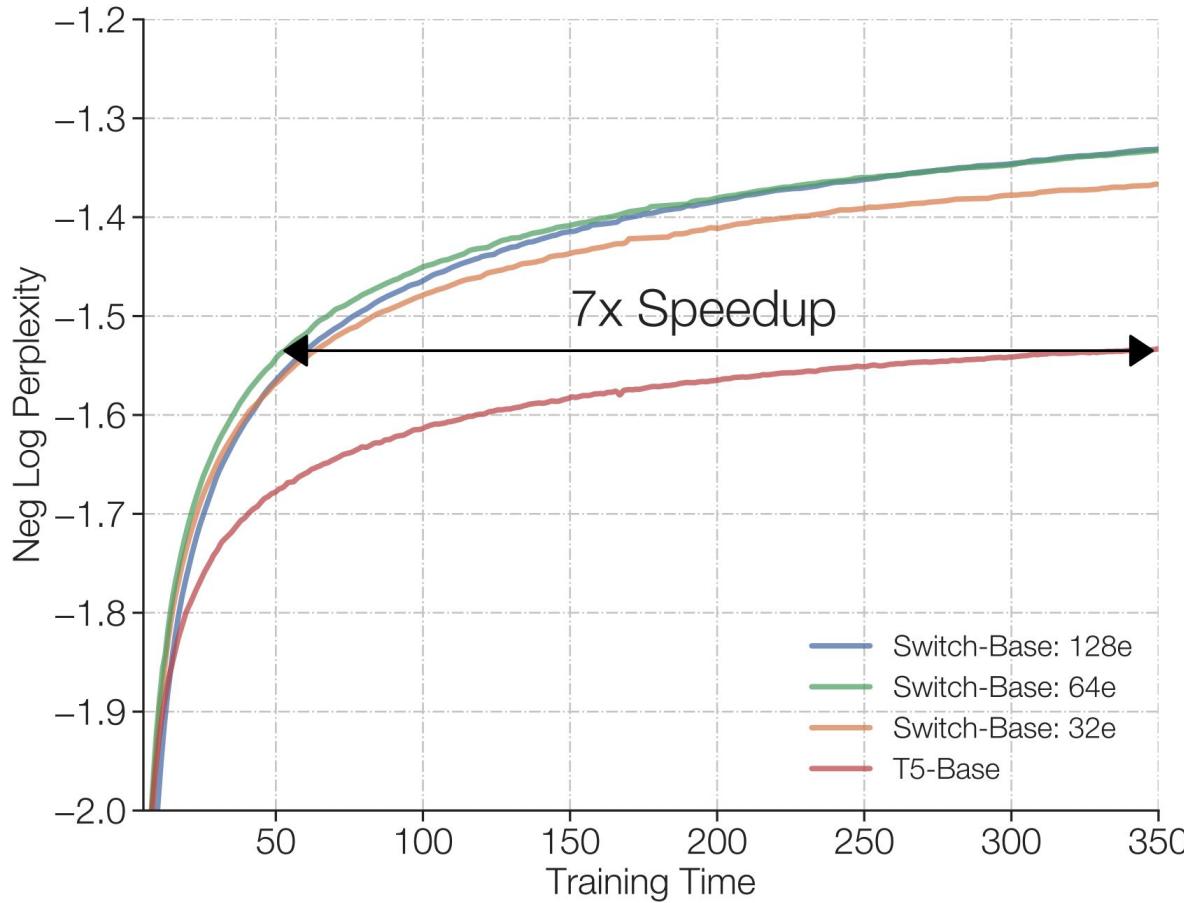
- The router assigns probabilities to each expert for every token.
- For each token, take the probability that expert  $i$  was preferred.
- Average over all tokens.

The auxiliary loss encourages uniform routing since it is minimized under a uniform distribution

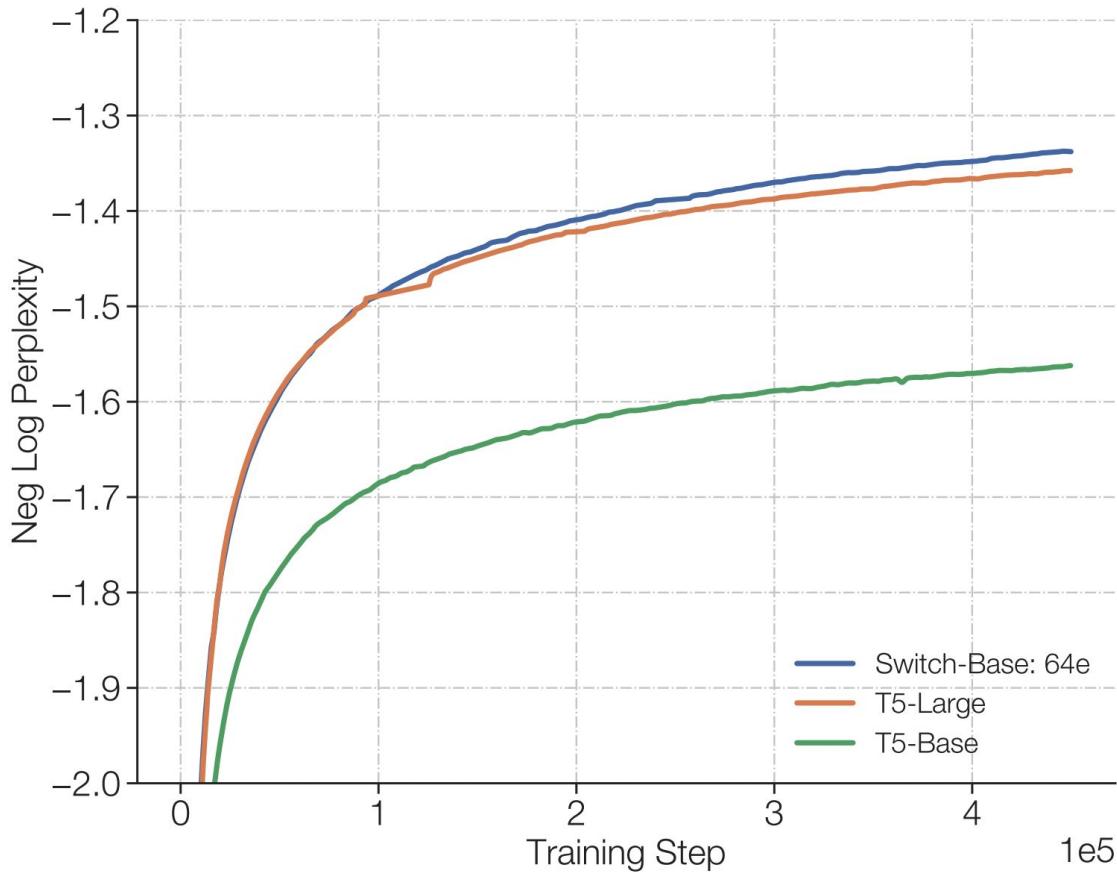
# Lower standard dropout rate for non-expert layers, higher for expert feed-forward layers

Model (dropout)	GLUE	CNNDM	SQuAD	SuperGLUE
T5-Base (d=0.1)	82.9	<b>19.6</b>	83.5	72.4
Switch-Base (d=0.1)	84.7	19.1	<b>83.7</b>	<b>73.0</b>
Switch-Base (d=0.2)	84.4	19.2	<b>83.9</b>	<b>73.2</b>
Switch-Base (d=0.3)	83.9	19.6	83.4	70.7
Switch-Base (d=0.1, ed=0.4)	<b>85.2</b>	<b>19.6</b>	<b>83.7</b>	<b>73.0</b>

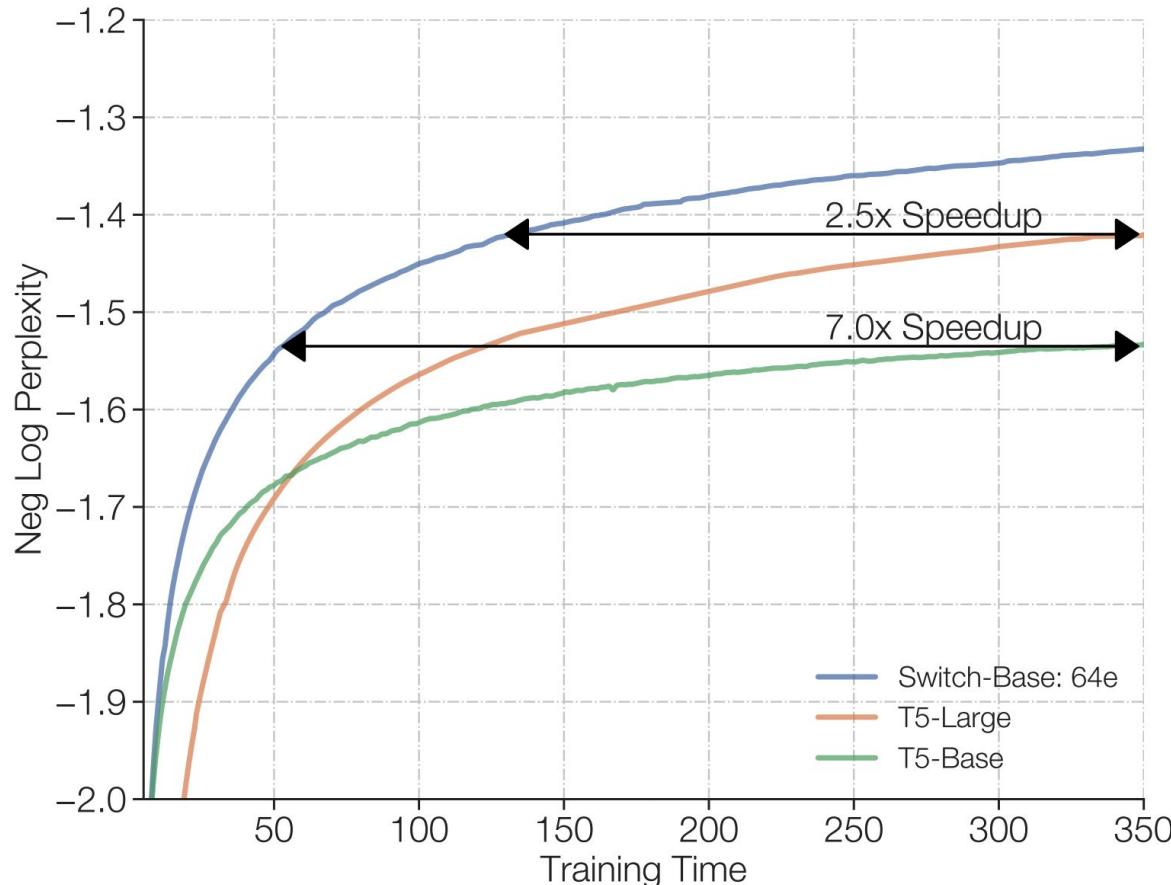
# Speed advantage of Switch Transformer



# Switch-Base is more sample efficient than T5-Large



# Switch-Base is faster than T5-Large (2.5x speedup)



# ... and significant downstream improvements

Model	GLUE	SQuAD	SuperGLUE	Winogrande (XL)
T5-Base	84.3	85.5	75.1	66.6
Switch-Base	<b>86.7</b>	<b>87.2</b>	<b>79.5</b>	<b>73.3</b>
T5-Large	87.8	88.1	82.7	79.1
Switch-Large	<b>88.5</b>	<b>88.6</b>	<b>84.7</b>	<b>83.0</b>

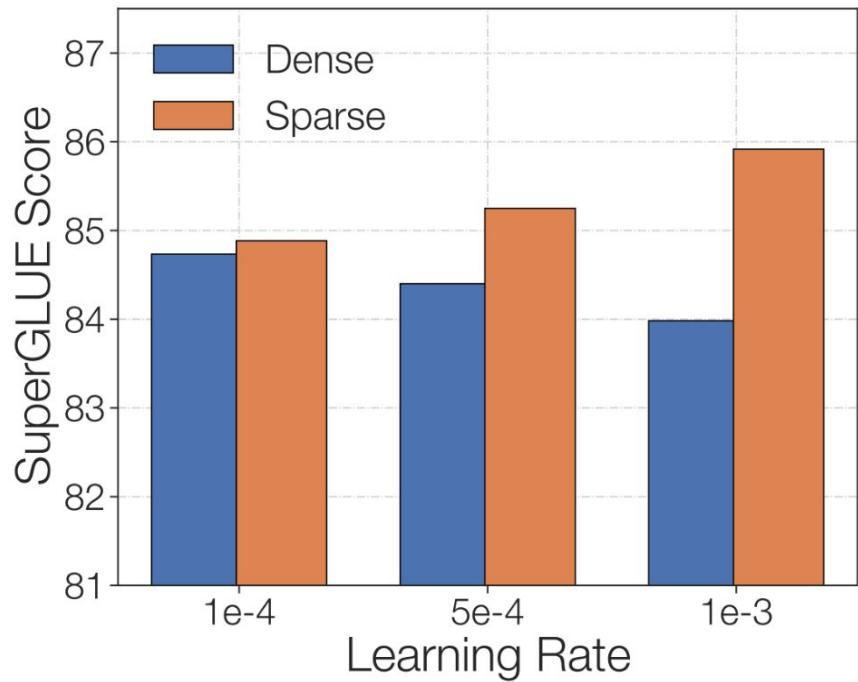
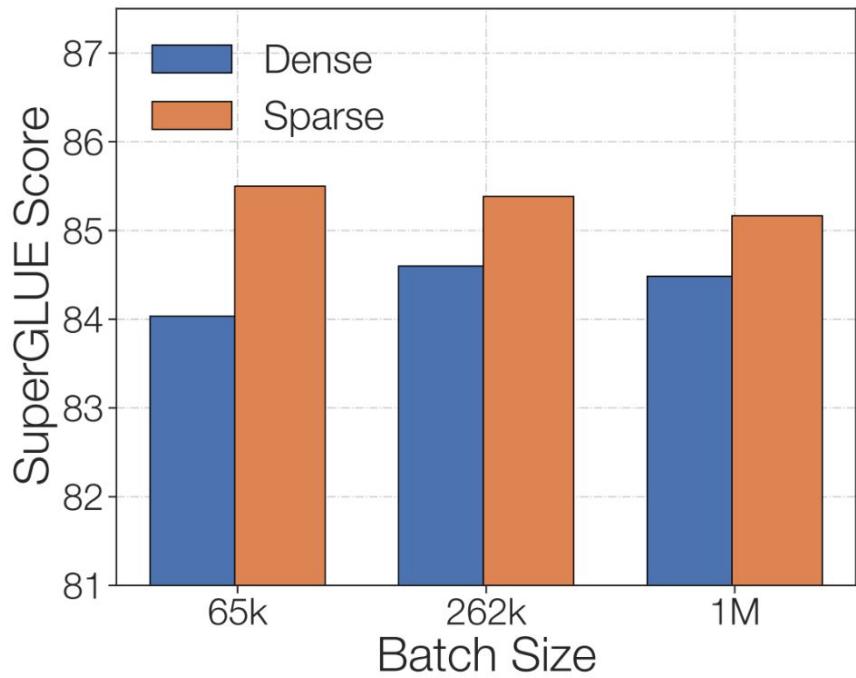
  

Model	XSum	ANLI (R3)	ARC Easy	ARC Chal.
T5-Base	18.7	51.8	56.7	<b>35.5</b>
Switch-Base	<b>20.3</b>	<b>54.0</b>	<b>61.3</b>	32.8
T5-Large	20.9	56.6	<b>68.8</b>	<b>35.5</b>
Switch-Large	<b>22.3</b>	<b>58.6</b>	66.0	<b>35.5</b>

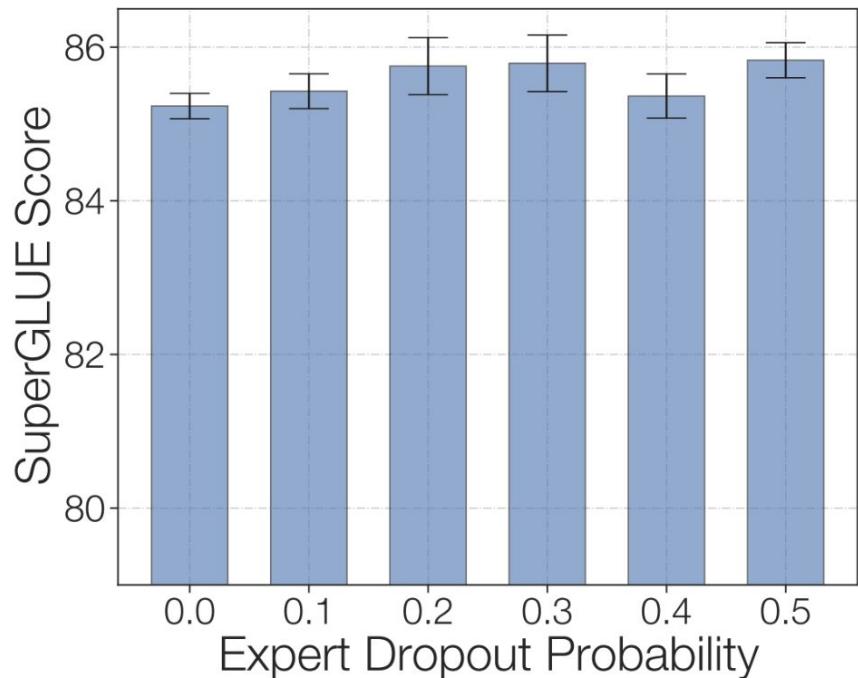
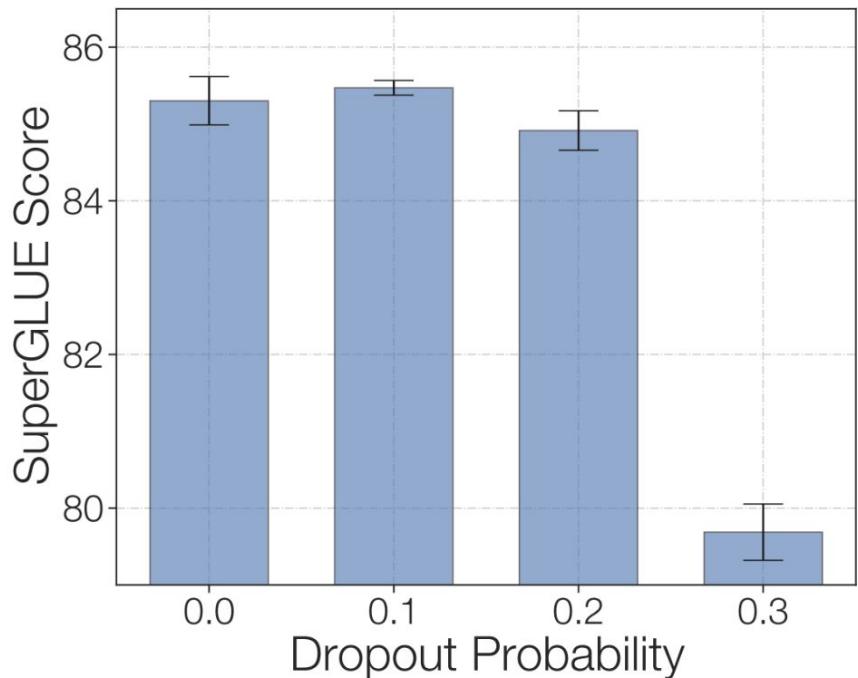
Model	CB Web QA	CB Natural QA	CB Trivia QA
T5-Base	26.6	25.8	24.5
Switch-Base	<b>27.4</b>	<b>26.8</b>	<b>30.7</b>
T5-Large	27.7	27.6	29.5
Switch-Large	<b>31.3</b>	<b>29.5</b>	<b>36.9</b>

# Sparse models benefit from small batch sizes and high learning rates



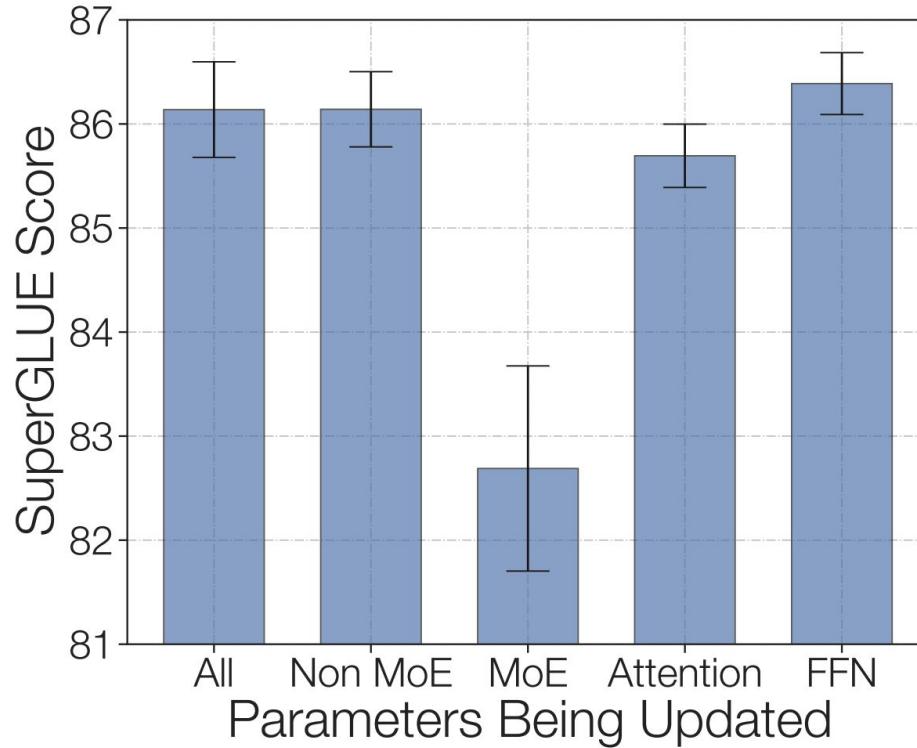
[“ST-MoE: Designing Stable and Transferable Sparse Expert Models” by Zoph et al. \(2022\)](#)

# Sparse models benefit from high dropout rates

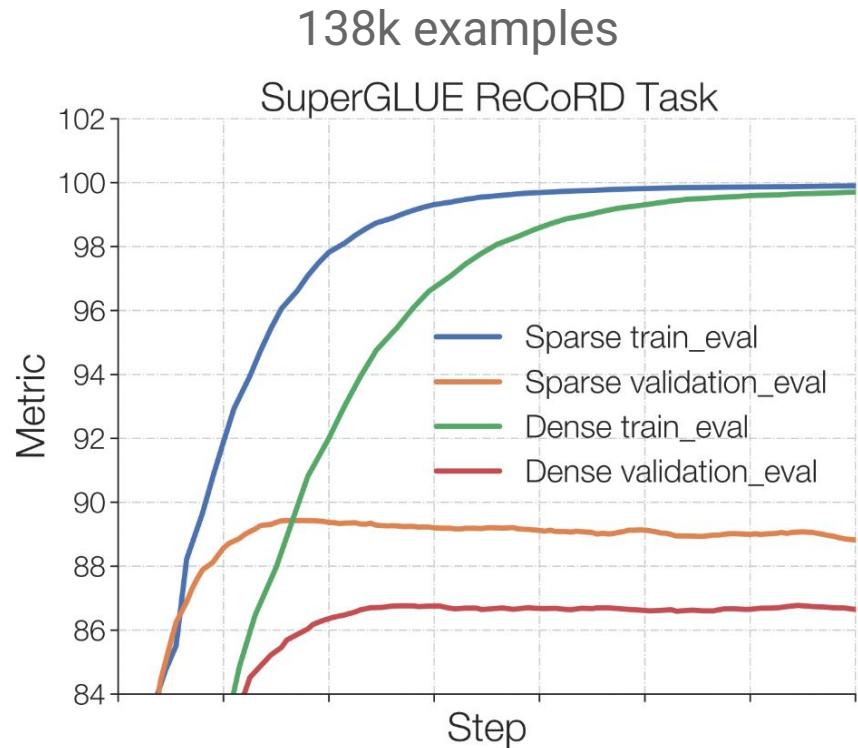
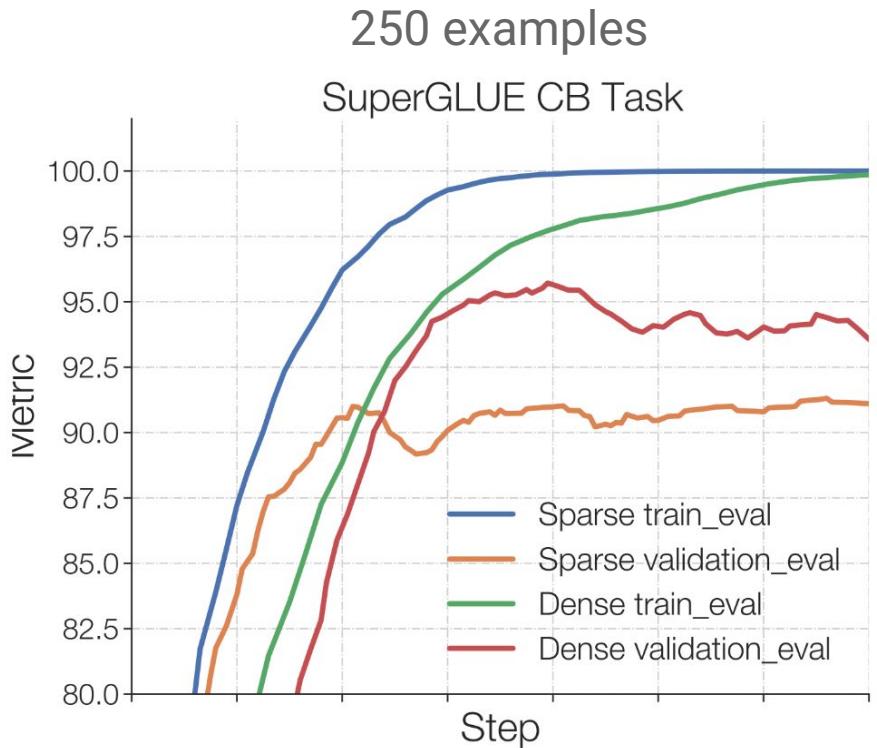


[“ST-MoE: Designing Stable and Transferable Sparse Expert Models” by Zoph et al. \(2022\)](#)

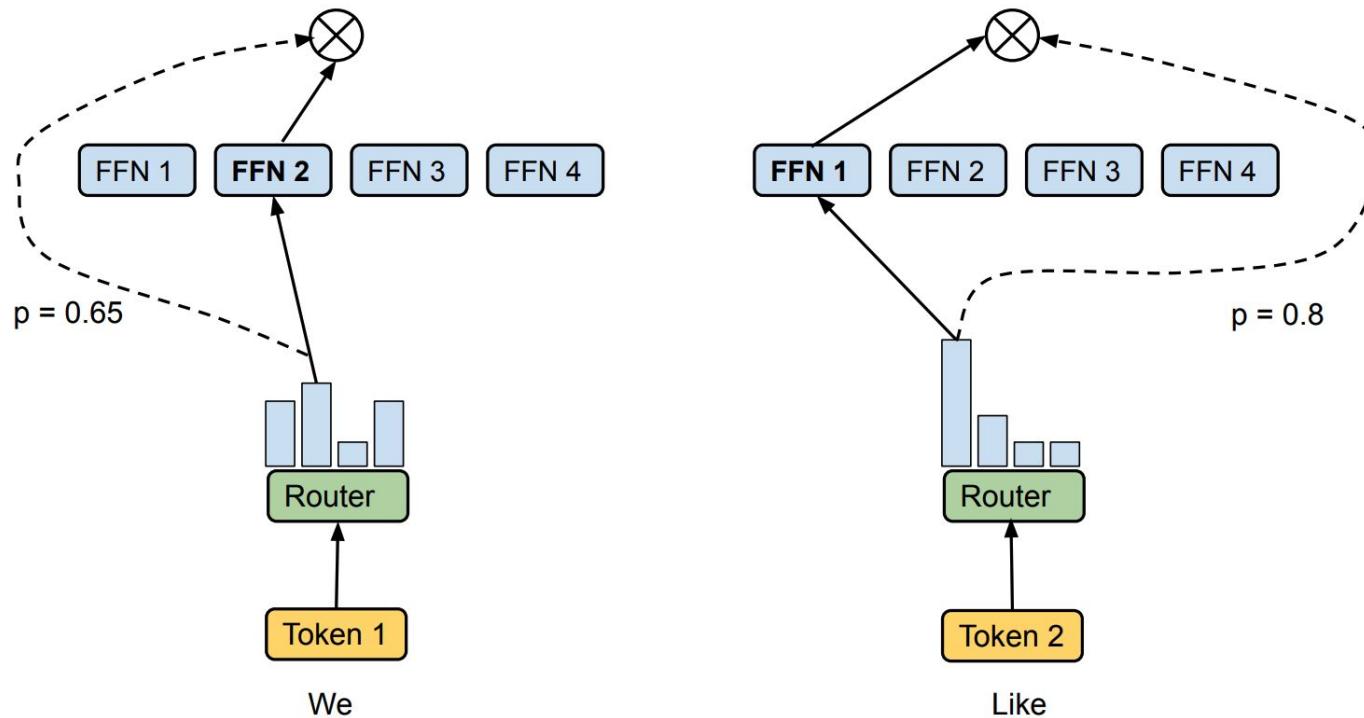
**For fine-tuning: by freezing the MoE layers, we can speed up the training while preserving the quality**



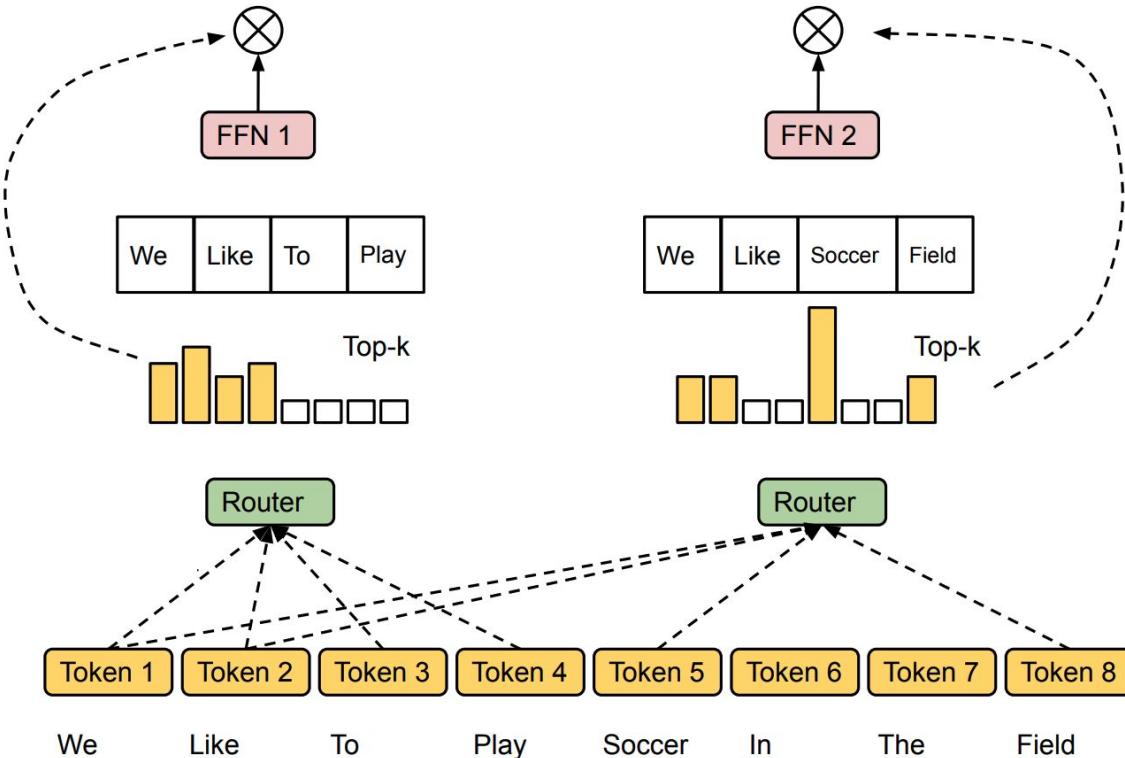
# Sparse models are prone to overfit



# Token-choice routing



# Expert-choice routing



["Mixture-of-Experts with Expert Choice Routing" by Zhou et al. \(2022\)](#)

---

# Mixture-of-Experts Meets Instruction Tuning: A Winning Combination for Large Language Models

---

**Sheng Shen<sup>¶\*</sup>**   **Le Hou<sup>†</sup>**   **Yanqi Zhou<sup>†</sup>**   **Nan Du<sup>†</sup>**   **Shayne Longpre<sup>⊤\*</sup>**   **Jason Wei<sup>†</sup>,**

**Hyung Won Chung<sup>†</sup>**   **Barret Zoph<sup>†</sup>**   **William Fedus<sup>†</sup>**   **Xinyun Chen<sup>†</sup>**   **Tu Vu<sup>‡\*,</sup>**

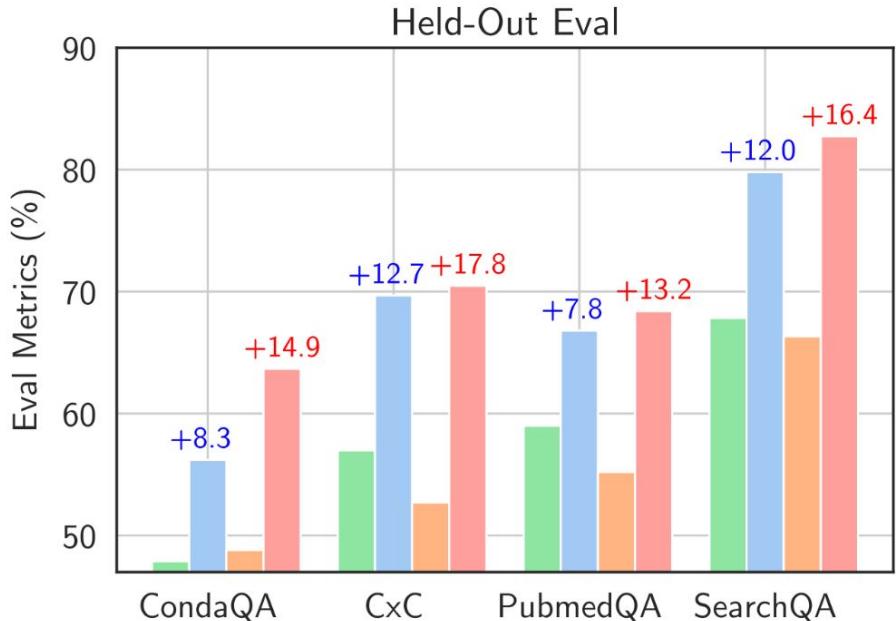
**Yuxin Wu<sup>†</sup>**   **Wuyang Chen<sup>§\*</sup>**   **Albert Webson<sup>†</sup>**   **Yunxuan Li<sup>†</sup>**   **Vincent Zhao<sup>†</sup>**   **Hongkun Yu<sup>†</sup>**

**Kurt Keutzer<sup>¶</sup>**   **Trevor Darrell<sup>¶</sup>**   **Denny Zhou<sup>†</sup>**

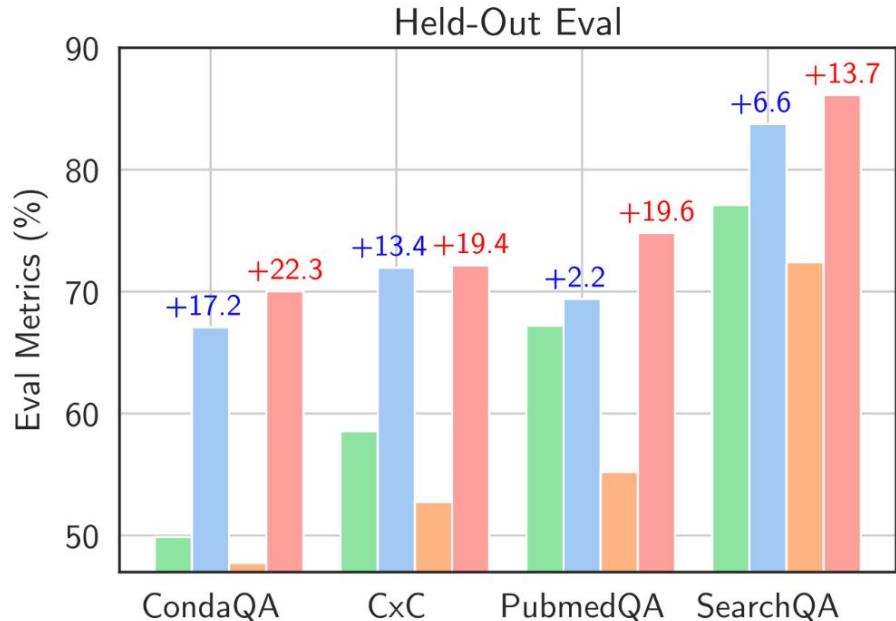
<sup>†</sup>Google    <sup>¶</sup>University of California, Berkeley    <sup>⊤</sup>Massachusetts Institute of Technology

<sup>‡</sup>University of Massachusetts Amherst    <sup>§</sup>The University of Texas at Austin

# Mixture-of-Experts meets Instruction Tuning



(a)  $\text{FLAN-EC}_{\text{BASE}}$  v.s.  $\text{FLAN-T5}_{\text{BASE}}$



(b)  $\text{FLAN-EC}_{\text{LARGE}}$  v.s.  $\text{FLAN-T5}_{\text{LARGE}}$



# When to use sparse MoEs vs dense models?

Experts are useful for high throughput scenarios with many machines. Given a fixed compute budget for pretraining, a sparse model will be more optimal. For low throughput scenarios with little VRAM, a dense model will be better.

**Thank you!**